



Building *Observable* Go Services

GopherCon Turkey - 2020



Who am I?

SWE/SRE **@Red Hat**
Observability Platform Team
Thanos Maintainer
Promethues Contributor

 kkakkoyun
 kakkojun





What is Observability?



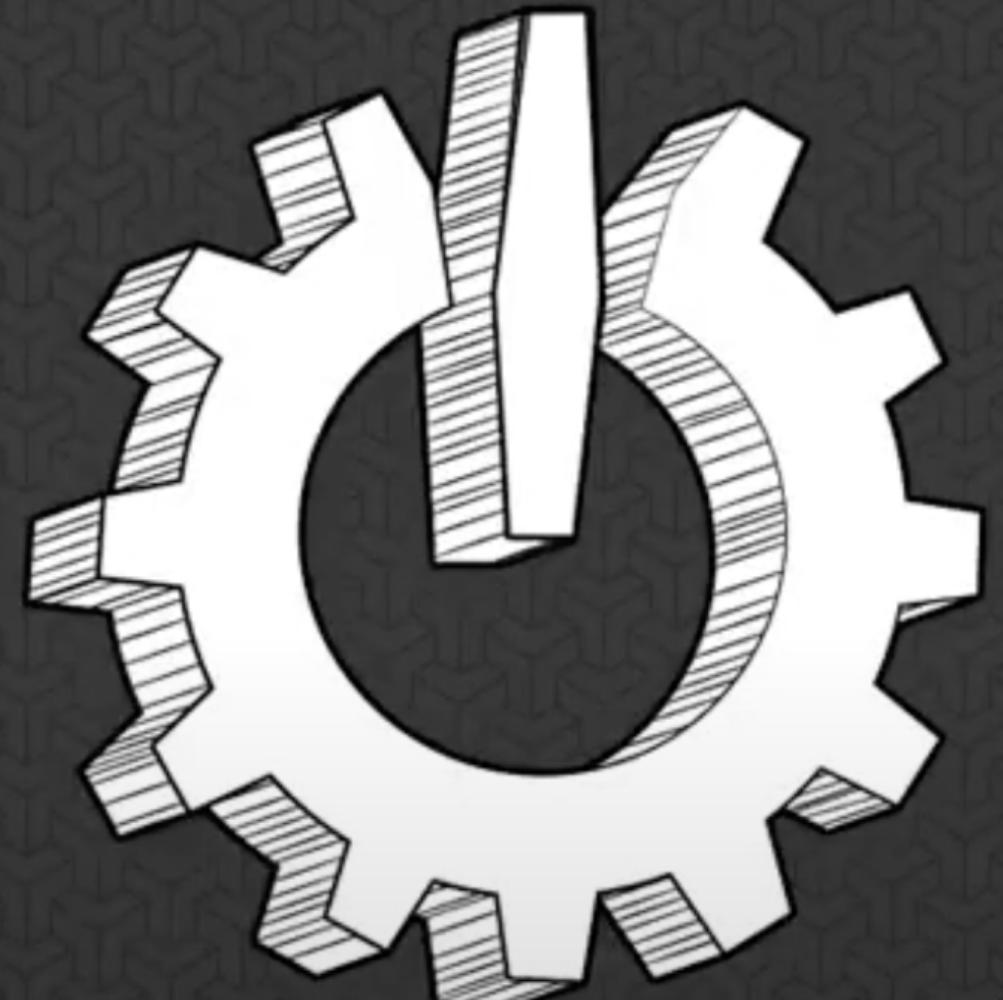
“In control theory, observability is a measure of how well internal states of a system can be inferred from knowledge of its external outputs.

- Wikipedia



Observability 101 ↗

FOSDEM 19
.org



On Observability

Observability 101

Richard Hartmann

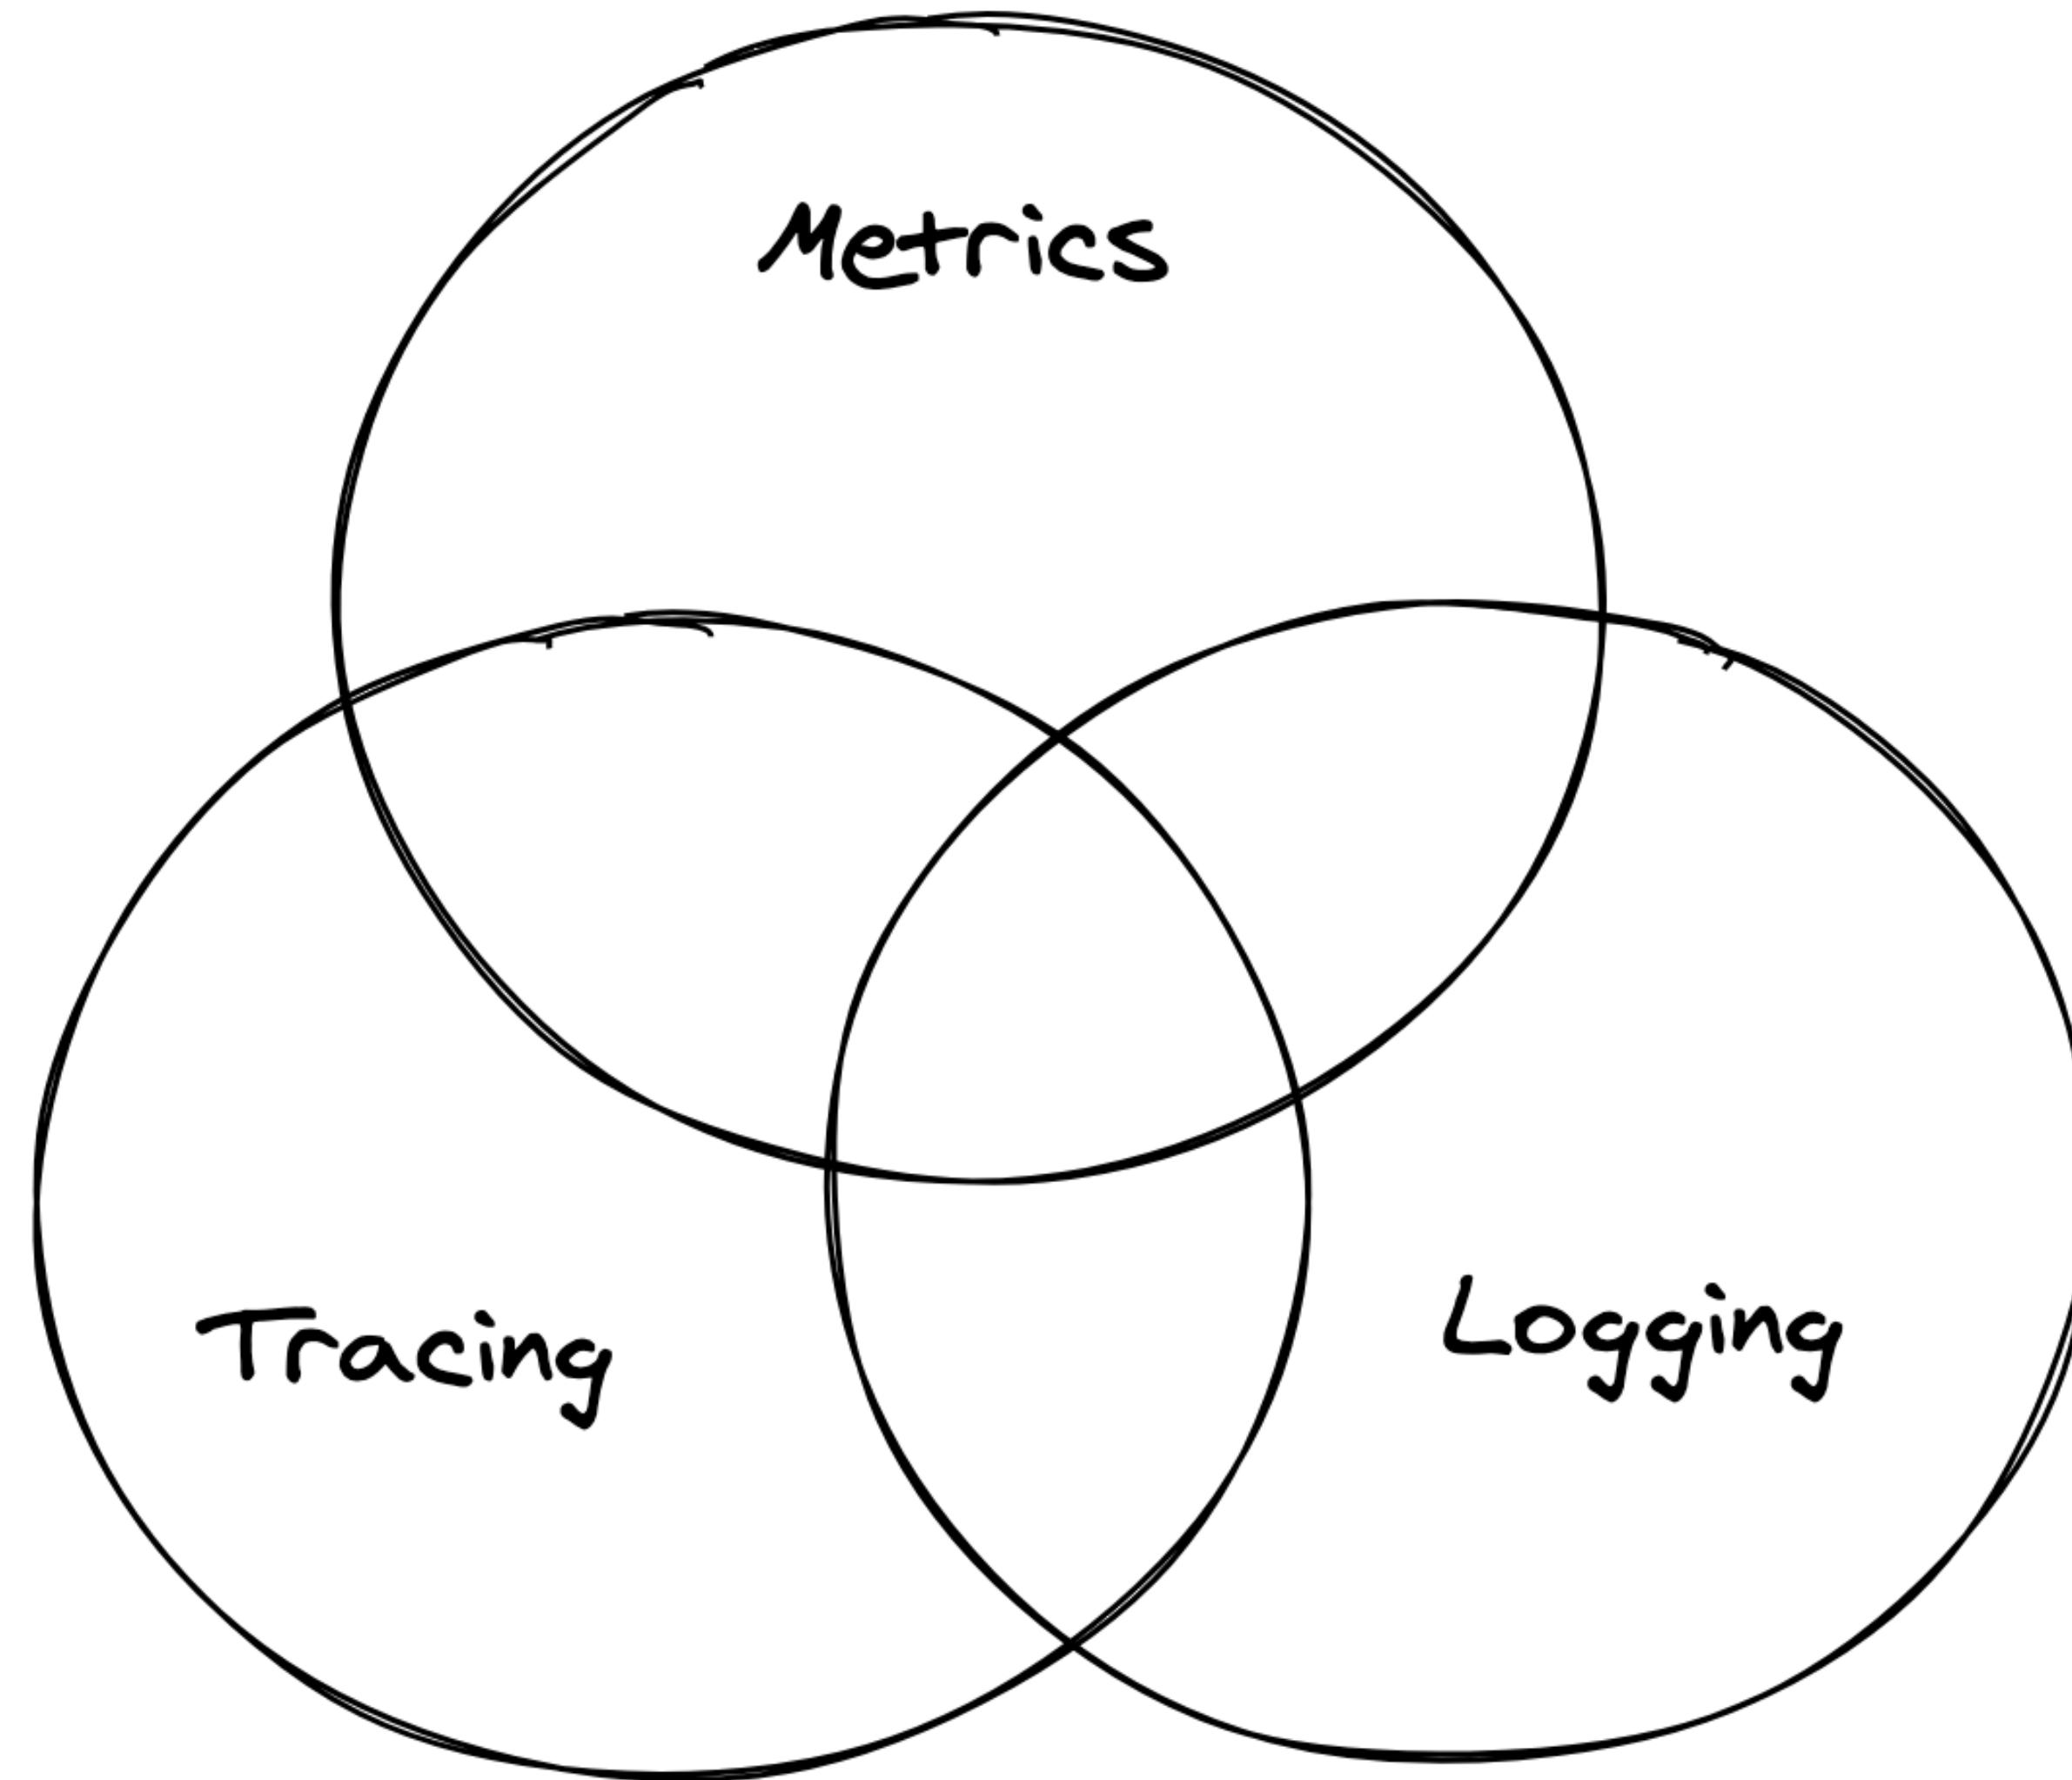
UB2.252A (Lameere)
2019-02-03



What does a properly
observable system in Go look
like?



The Pillars of Observability

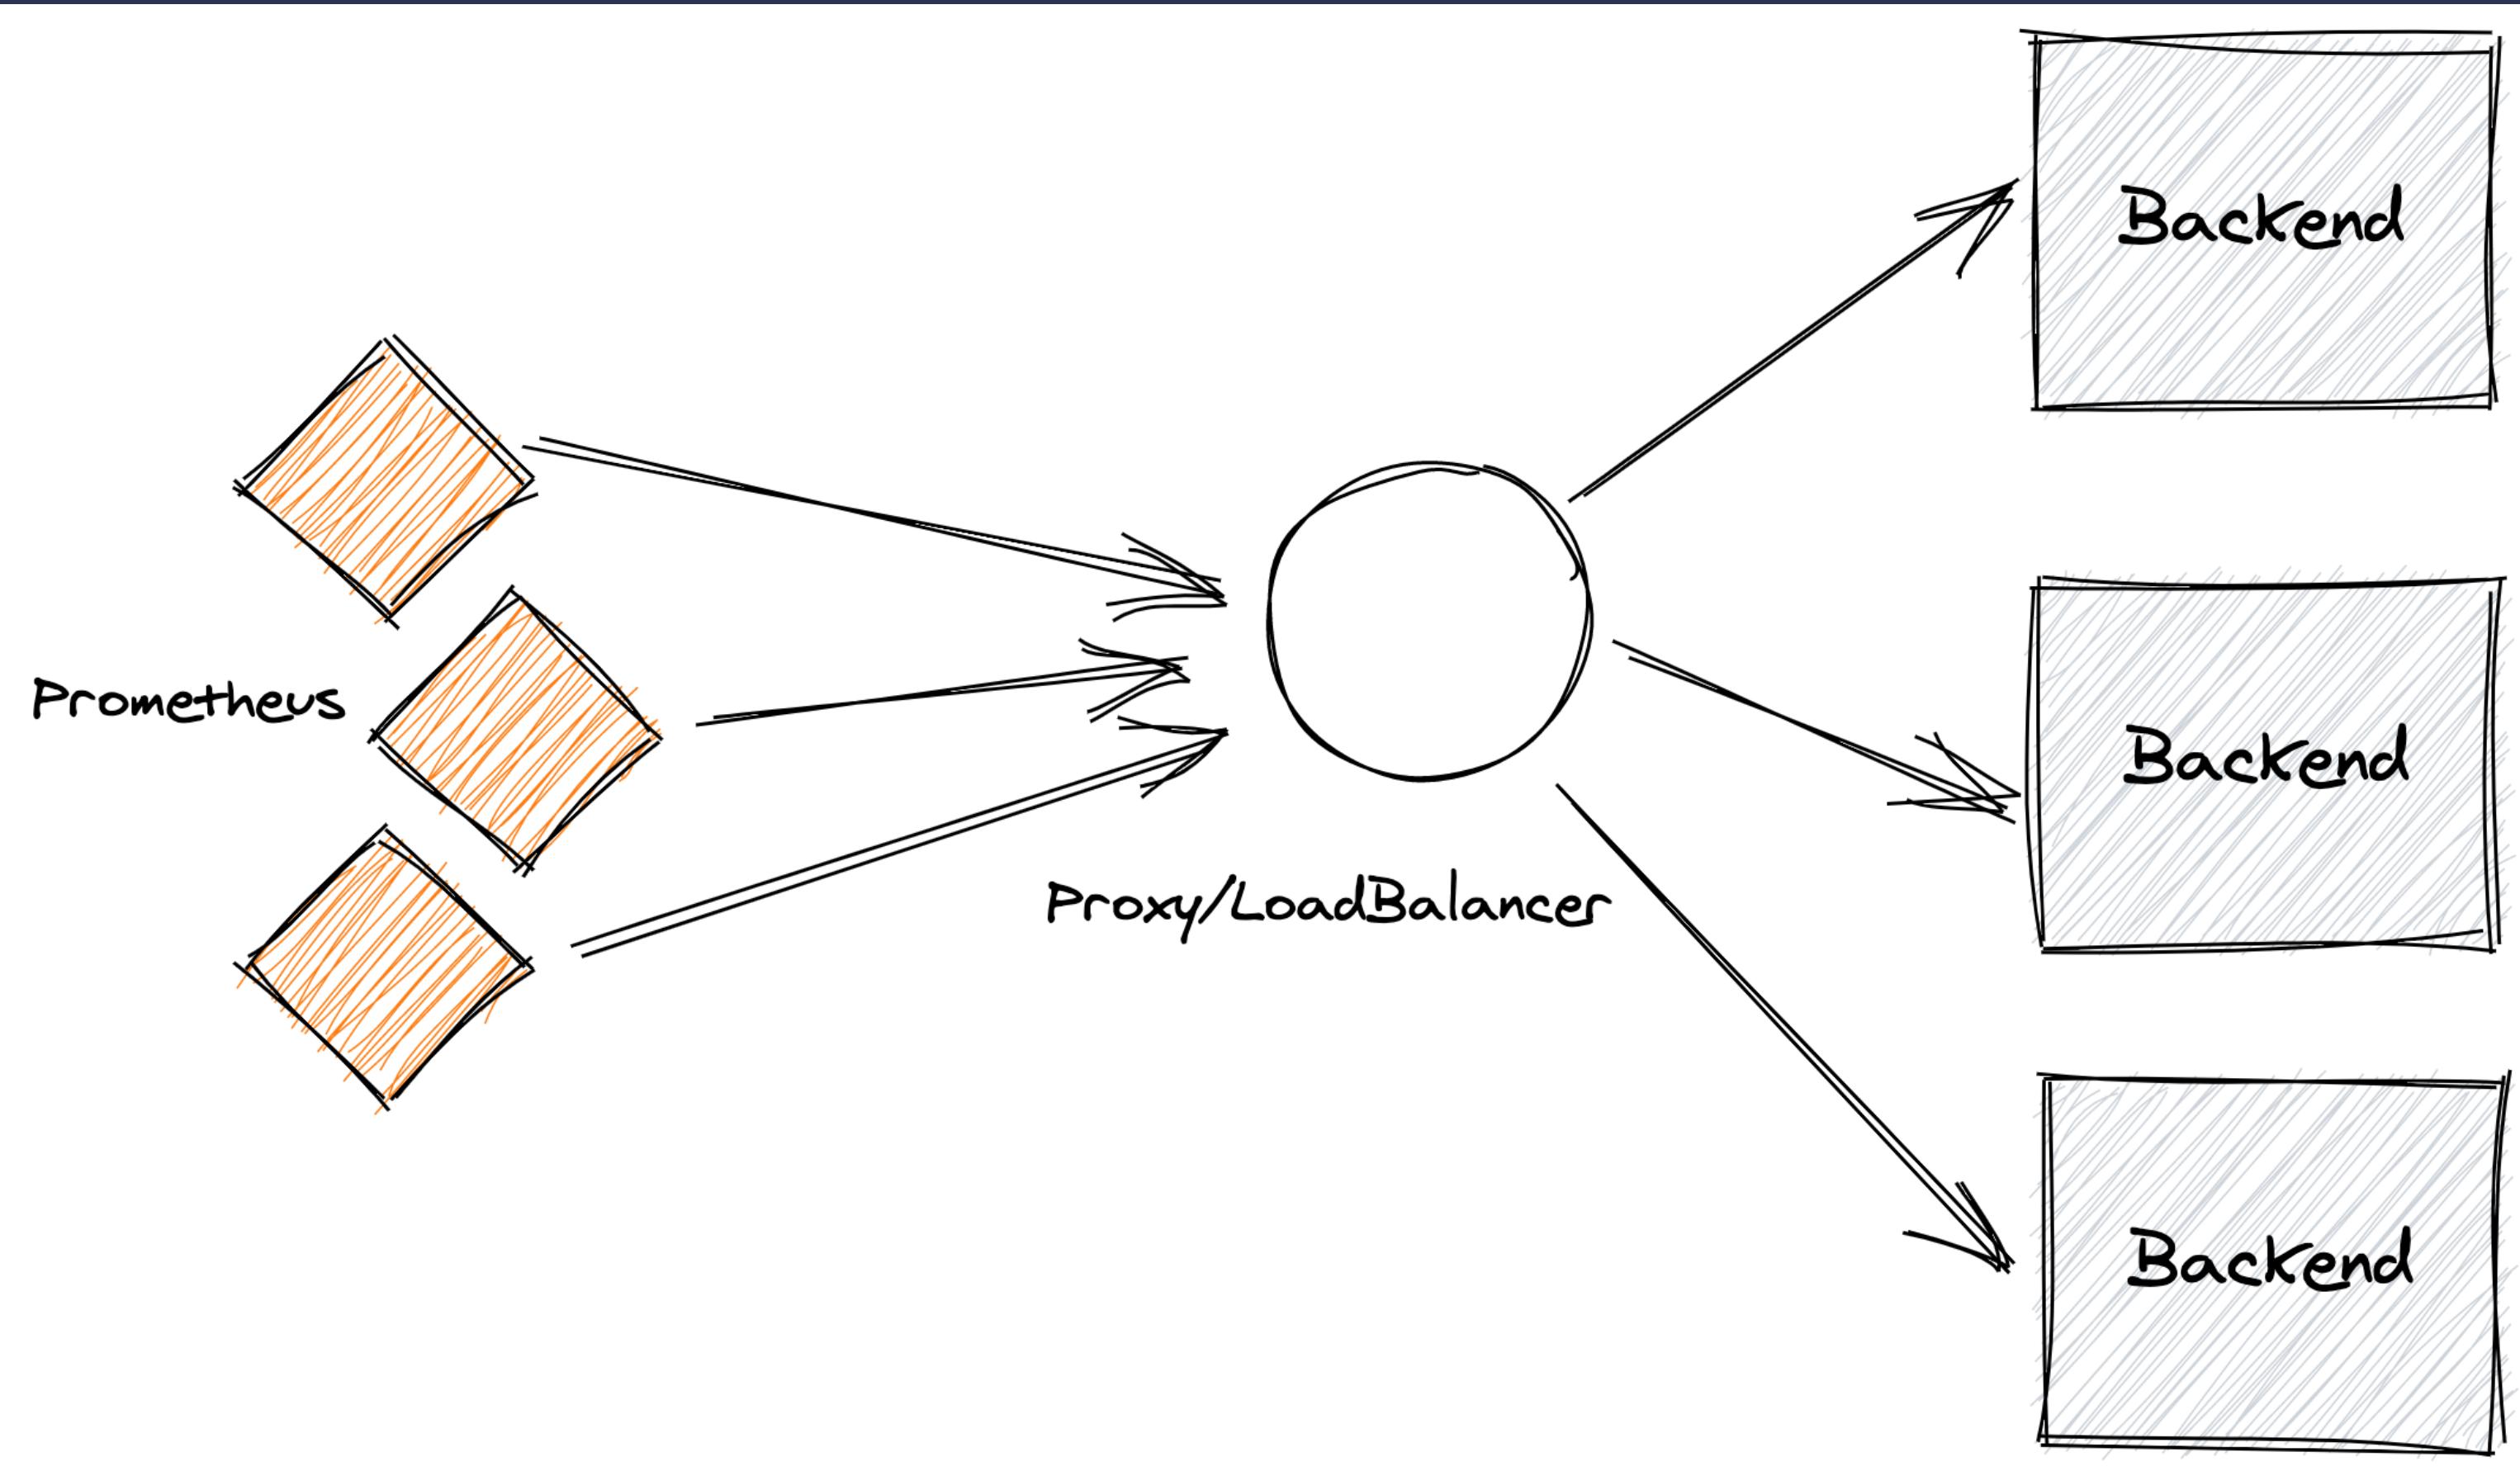




Metrics, tracing, and logging



Everything starts with Instrumentation





Health checks



```
// NewServer creates a new internal server that exposes debug probes.
func NewServer(reg prometheus.Gatherer, listen, healthcheckURL string) *l
    // Internal server to expose introspection APIs.
    mux := http.NewServeMux()

    // Initialize health checks.
    healthchecks := healthcheck.NewHandler()

    // Register health check endpoints.
    mux.Handle("/-/healthy", http.HandlerFunc(healthchecks.LiveEndpoint))
    mux.Handle("/-/ready", http.HandlerFunc(healthchecks.ReadyEndpoint))

    return &http.Server{
        Addr:   listen,
        Handler: mux,
    }
}
```



```
☒ curl localhost:8081/-/heathy
{
    "http": "OK"
}

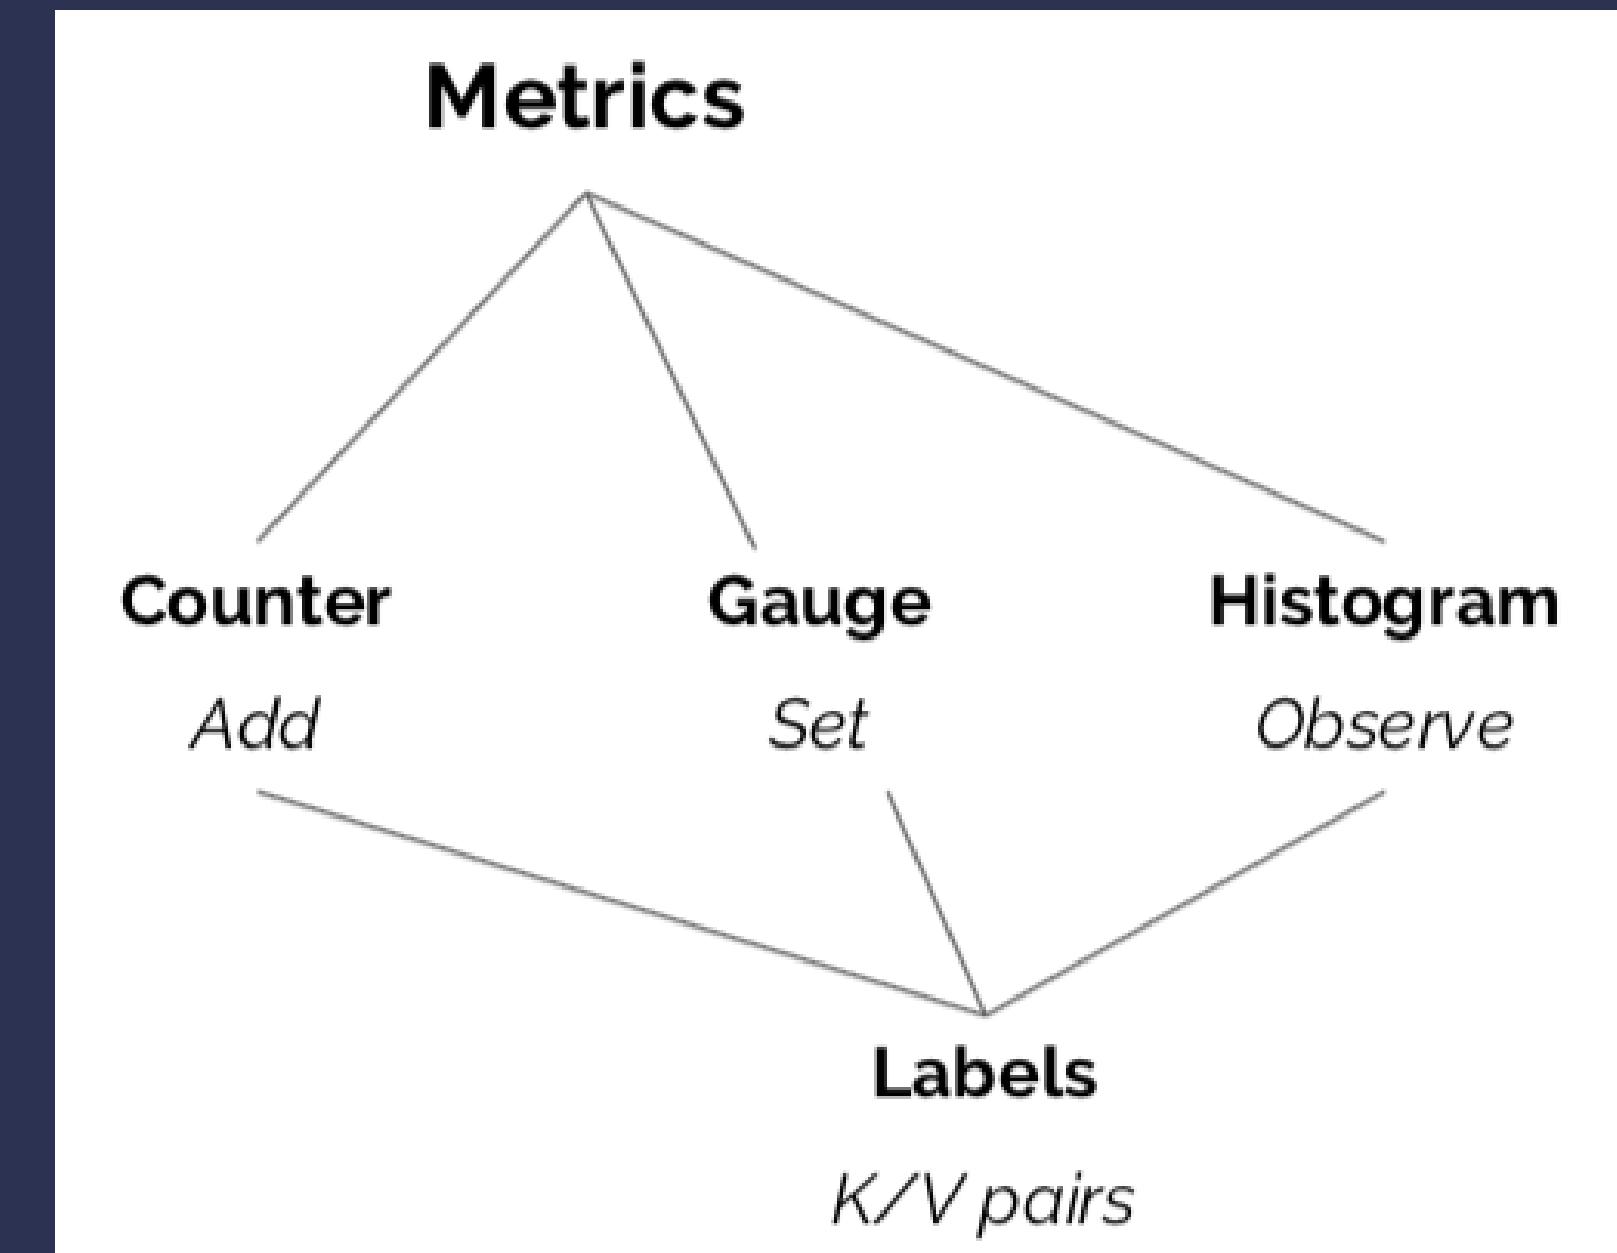
☒ curl localhost:8081/-/ready
{
    "http": "OK"
}
```



metalmatze/signal 
heptio/healthcheck 



Metrics





Prometheus

[DOCS](#)[DOWNLOAD](#)[COMMUNITY](#)[SUPPORT & TRAINING](#)[BLOG](#)

From metrics to insight

Power your metrics and alerting with a leading open-source monitoring solution.

[GET STARTED](#)[DOWNLOAD](#)

Dimensional data

Prometheus implements a highly dimensional data model. Time series are identified by a metric name and a set of key-value pairs.

Powerful queries

PromQL allows slicing and dicing of collected time series data in order to generate ad-hoc graphs, tables, and alerts.

Great visualization

Prometheus has multiple modes for visualizing data: a built-in expression browser, Grafana integration, and a console template language.

Efficient storage

Prometheus stores time series in memory and on local disk in an efficient custom format. Scaling is achieved by functional sharding and federation.

Simple operation

Each server is independent for reliability, relying only on local storage.

Precise alerting

Alerts are defined based on Prometheus's flexible PromQL and

Many client libraries

Client libraries allow easy instrumentation of services. Over ten

Many integrations

Existing exporters allow bridging of third-party data into Prometheus.



[prometheus/client_golang](#) ↗



```
package main

import "github.com/prometheus/client_golang/prometheus"

func main() {
    // Start metric registry.
    reg := prometheus.NewRegistry()

    // Register metrics server.
    http.Handle("/metrics", promhttp.HandlerFor(reg, promhttp.HandlerOpts))
    http.ListenAndServe(":8080", nil)
}
```



```
✉ curl localhost:8081/metrics
# HELP go_gc_duration_seconds A summary of the pause duration of garbage
# TYPE go_gc_duration_seconds summary
go_gc_duration_seconds{quantile="0"} 0
go_gc_duration_seconds{quantile="0.25"} 0
go_gc_duration_seconds{quantile="0.5"} 0
go_gc_duration_seconds{quantile="0.75"} 0
go_gc_duration_seconds{quantile="1"} 0
go_gc_duration_seconds_sum 0
go_gc_duration_seconds_count 0
...
...
```



```
import (
    "github.com/povilasv/prommod"
    "github.com/prometheus/common/version"
)

serviceName = "observable_remote_write_backend"

func main() {
    reg := prometheus.NewRegistry()

    // Register standard Go metric collectors, which are by default registered
    reg.MustRegister(
        prometheus.NewGoCollector(),
        version.NewCollector(serviceName),
        prommod.NewCollector(serviceName),
        prometheus.NewProcessCollector(prometheus.ProcessCollectorOpts{})
    )
}
```



Enable query history

build_info

`observable_remote_write_backend_build_info`

`observable_remote_write_proxy_build_info`



Moment



Element

no data

Add Graph



Prometheus Alerts Graph Status ▾ Help

Enable query history

go_

- go_gc_duration_seconds
- go_gc_duration_seconds_count
- go_gc_duration_seconds_sum
- go_goroutines
- go_info
- go_memstats_alloc_bytes
- go_memstats_alloc_bytes_total
- go_memstats_buck_hash_sys_bytes
- go_memstats_frees_total
- go_memstats_gc_cpu_fraction
- go_memstats_gc_sys_bytes
- go_memstats_heap_alloc_bytes
- go_memstats_heap_idle_bytes
- go_memstats_heap_inuse_bytes
- go_memstats_heap_objects
- go_memstats_heap_released_bytes
- go_memstats_heap_sys_bytes
- go_memstats_last_gc_time_seconds
- go_memstats_lookups_total
- go_memstats_mallocs_total
- go_memstats_mcache_inuse_bytes

 Enable query history[Try experimental React UI](#)

go_mod_info



Load time: 45ms
Resolution: 14s
Total time series: 93

[Execute](#)

- insert metric at cursor -

[Graph](#)[Console](#)

Moment

**Element****Value**

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/apache/thrift",program="observable_remote_write_backend",version="v0.13.0"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/beorn7/perks",program="observable_remote_write_backend",version="v1.0.1"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/cespare/xxhash/v2",program="observable_remote_write_backend",version="v2.1.1"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/go-chi/chi",program="observable_remote_write_backend",version="v4.1.2+incompatible"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/go-kit/kit",program="observable_remote_write_backend",version="v0.10.0"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/go-logfmt/logfmt",program="observable_remote_write_backend",version="v0.5.0"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/gogo/protobuf",program="observable_remote_write_backend",version="v1.3.1"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/golang/protobuf",program="observable_remote_write_backend",version="v1.4.2"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/golang/snappy",program="observable_remote_write_backend",version="v0.0.1"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/grpc-ecosystem/grpc-gateway",program="observable_remote_write_backend",version="v1.14.6"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/matttproud/golang_protobuf_extensions",program="observable_remote_write_backend",version="v1.0.1"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/metalmatze/signal",program="observable_remote_write_backend",version="v0.0.0-20200616171423-be84551ba3ce"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/oklog/run",program="observable_remote_write_backend",version="v1.1.0"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/oklog/ulid",program="observable_remote_write_backend",version="v1.3.1"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/pkg/errors",program="observable_remote_write_backend",version="v0.9.1"}

1

go_mod_info{instance="127.0.0.1:8180",job="backend",name="github.com/povilasv/prommod",program="observable_remote_write_backend",version="v0.0.12"}

1



```
package middleware

import (
    "github.com/prometheus/client_golang/prometheus"
)

type MetricsMiddleware struct {
    requestDuration *prometheus.HistogramVec
    requestSize     *prometheus.SummaryVec
    requestsTotal   *prometheus.CounterVec
    responseSize    *prometheus.SummaryVec
}
```



```
requestDuration: promauto.With(reg).NewHistogramVec(  
    prometheus.HistogramOpts{  
        Name:      "http_request_duration_seconds",  
        Help:      "Tracks the latencies for HTTP requests.",  
        Buckets:  []float64{0.001, 0.01, 0.1, 0.3, 0.6, 1, 3, 6, 9, 20},  
        Labels:   []string{"code", "handler", "method"},  
    },  
  
    requestSize: promauto.With(reg).NewSummaryVec(  
        prometheus.SummaryOpts{  
            Name:      "http_request_size_bytes",  
            Help:      "Tracks the size of HTTP requests.",  
            Labels:   []string{"code", "handler", "method"},  
        },  
    ),
```



```
import (
    "github.com/prometheus/client_golang/prometheus"
    "github.com/prometheus/client_golang/prometheus/promauto"
    "github.com/prometheus/client_golang/prometheus/promhttp"
)

func (ins *MetricsMiddleware) NewHandler(handlerName string) func(next ht
    return func(handler http.Handler) http.Handler {
        return promhttp.InstrumentHandlerDuration(
            ins.requestDuration.MustCurryWith(prometheus.Labels{"handler"})
        )
        return promhttp.InstrumentHandlerRequestSize(
            ins.requestSize.MustCurryWith(prometheus.Labels{"handler"})
        )
        return promhttp.InstrumentHandlerResponseSize(
            ins.responseSize.MustCurryWith(prometheus.Labels{
                "handler",
            }),
        ),
    },
)
```



```
metrics := middleware.NewMetricsMiddleware(reg)

// Main server to listen for public APIs.
mux := http.NewServeMux()
mux.Handle("/receive",
    metrics.NewHandler("receive")(receiver.Receive),
)

srv := &http.Server{
    Addr:    cfg.server.listen,
    Handler: mux,
}

srv.ListenAndServe()
```



Prometheus Alerts Graph Status ▾ Help

Enable query history

Try experimental React UI

```
histogram_quantile(0.50, sum by (le) (rate(http_request_duration_seconds_bucket{job="observatorium-thanos-query",code!~"5.",handler="query_range"}[1d])))
```



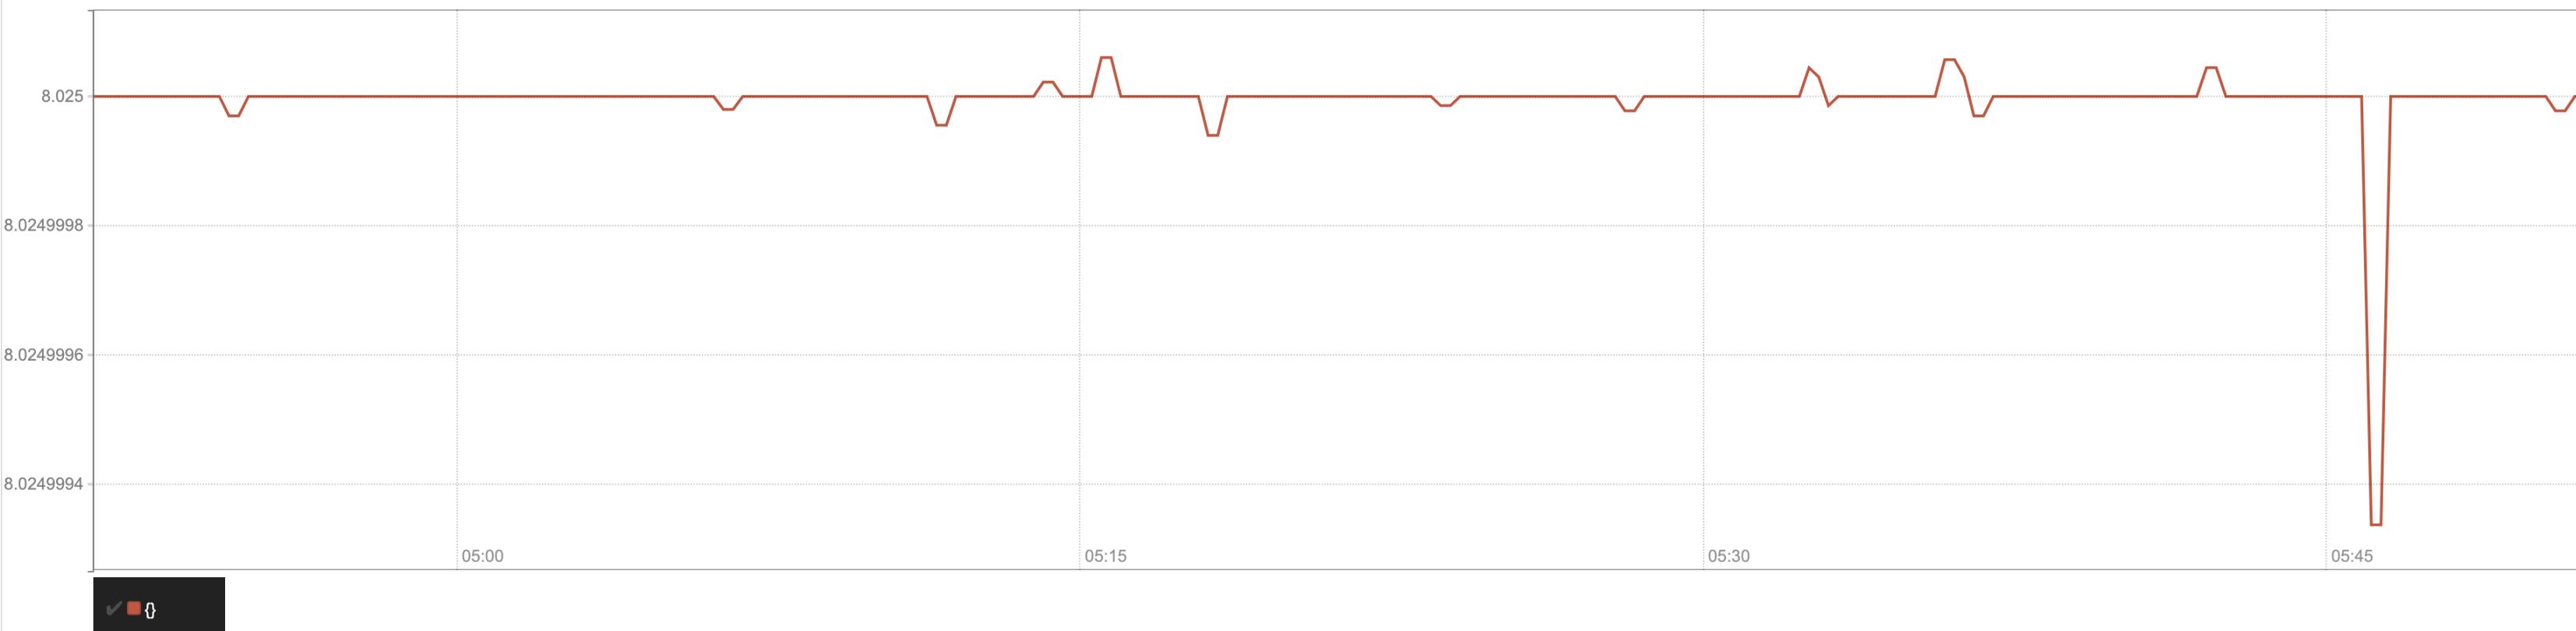
Execute

- insert metric at cursor ▾

Graph

Console

- 1h + ⏪ Until ⏩ Res. (s) □ stacked



Load time: 488ms
Resolution: 14s
Total time series: 1

Add Graph

Remove Graph



Logs



Structured {"db_err_code": 14002, ...}

Unstructured "[WARN] DB error: 14002 ..."

Logging



go-kit/logger ↗
uber/zap ↗



```
func NewLogger(logLevel, logFormat, debugName string) log.Logger {
    var (
        logger log.Logger
        lvl    level.Option
    )

    switch logLevel {
    case "error":
        lvl = level.AllowError()
    case "warn":
        lvl = level.AllowWarn()
    case "info":
        lvl = level.AllowInfo()
    case "debug":
        lvl = level.AllowDebug()
    default:
        panic("unexpected log level")
    }
    logger = log.New(lvl, logFormat, logName(debugName))
    return logger
}
```



```
func main() {
    // Initialize structured logger.
    logger := internal.NewLogger(
        cfg.logLevel,
        cfg.logFormat,
        cfg.debug.name,
    )
    defer level.Info(logger).Log("msg", "existing")
}
```



```
level.Warn(logger).Log("msg", "http read", "err", err)
level.Warn(logger).Log("msg", "snappy decode", "err", err)
level.Warn(logger).Log("msg", "proto unmarshalling", "err", err)
level.Info(logger).Log("msg", "remote write request received")
level.Debug(logger).Log("msg", m)
level.Debug(logger).Log("msg", fmt.Sprintf(" %f %d", s.Value, s.Timestamp))
```



```
level=info name=observable-remote-write-backend ts=2020-07-24T17:44:37Z
caller=main.go:127 msg="starting server"

level=info name=observable-remote-write-backend ts=2020-07-24T17:44:53Z
caller=receiver.go:64 msg="remote write request received"
```



```
// Logger returns a middleware to log HTTP requests.
func Logger(logger log.Logger) func(next http.Handler) http.Handler {
    return func(next http.Handler) http.Handler {
        return http.HandlerFunc(func(w http.ResponseWriter, r *http.Request) {
            start := time.Now()

            next.ServeHTTP(w, r)
            ...

            keyvals := []interface{}{
                "request", RequestIDFromContext(r.Context()),
                "proto", r.Proto,
                "method", r.Method,
                "path", r.URL.Path,
                "duration", time.Since(start),
            }

            level.Debug(logger).Log(keyvals...)
        })
    }
}
```



```
metrics := middleware.NewMetricsMiddleware(reg)

// Main server to listen for public APIs.
mux := http.NewServeMux()
mux.Handle("/receive",
    middleware.RequestID(
        middleware.Logger(logger)(receiver.Receive),
    ),
)

srv := &http.Server{
    Addr:    cfg.server.listen,
    Handler: mux,
}

srv.ListenAndServe()
```



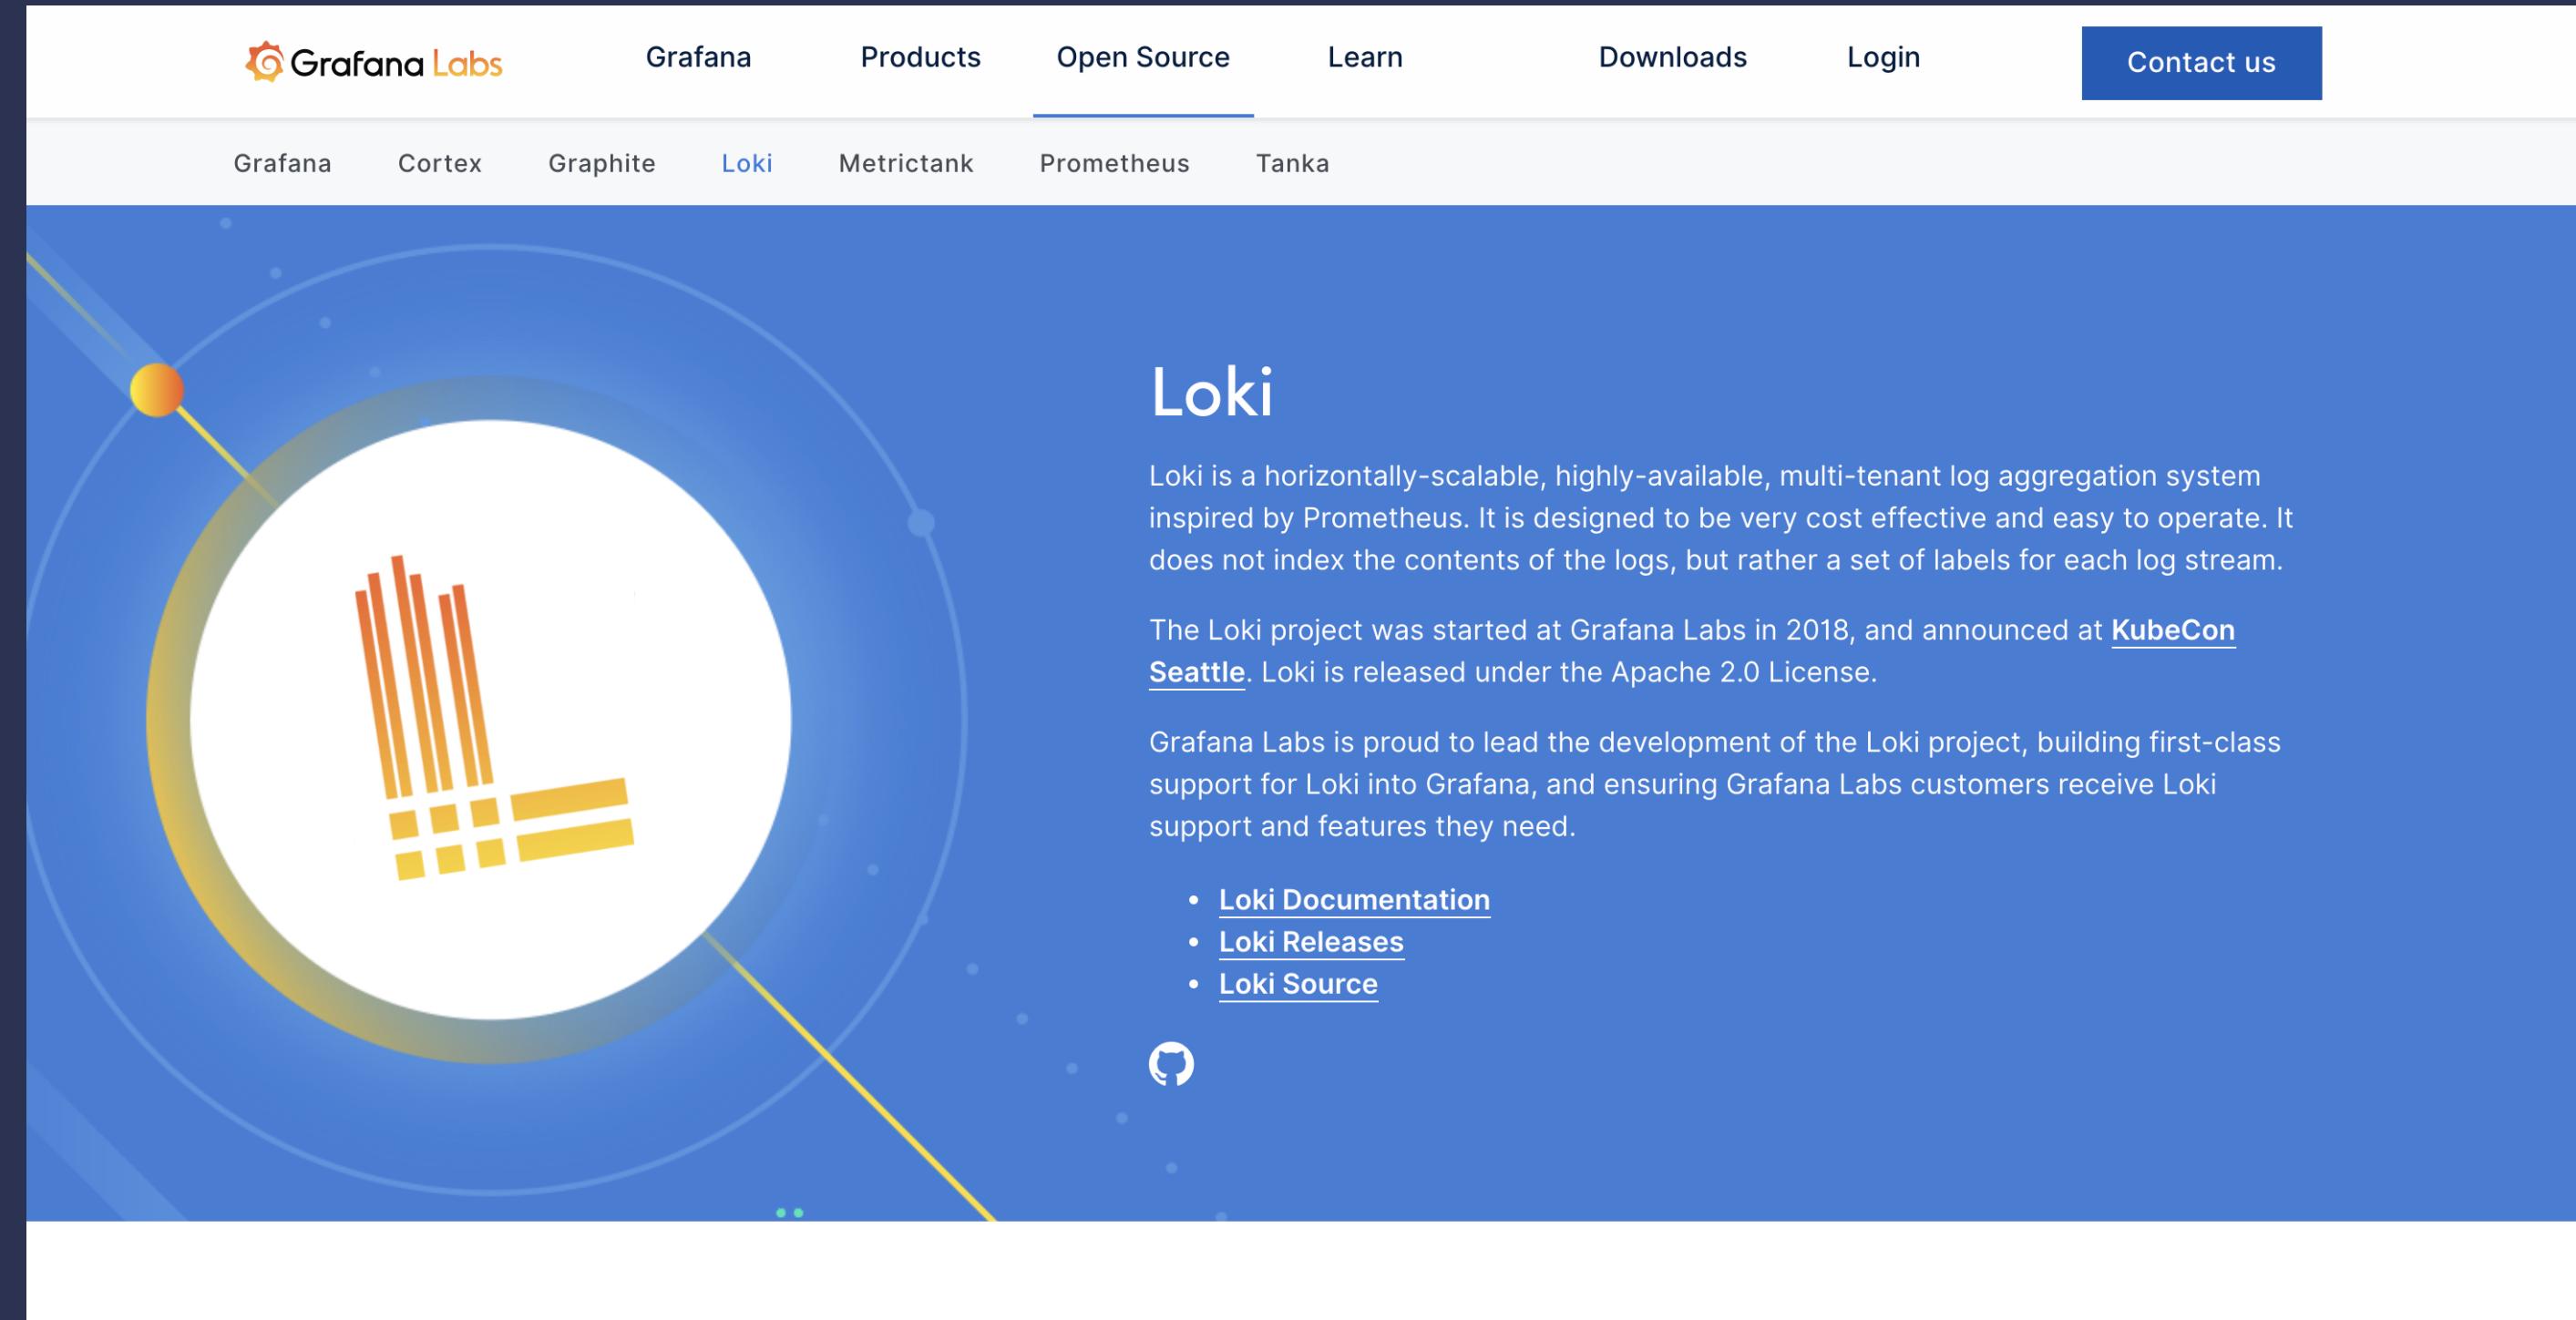
```
// RequestID returns a middleware that sets a unique request id for each
func RequestID(h http.Handler) http.Handler {
    return http.HandlerFunc(func(w http.ResponseWriter, r *http.Request)
        reqID := r.Header.Get("X-Request-ID")
        if reqID == "" {
            entropy := ulid.Monotonic(rand.New(rand.NewSource(time.Now())))
            reqID := ulid.MustNew(ulid.Timestamp(time.Now()), entropy)
            r.Header.Set("X-Request-ID", reqID.String())
        }
        ctx := newContextWithRequestID(r.Context(), reqID)
        h.ServeHTTP(w, r.WithContext(ctx))
    )
}
```



loki

[Grafana Labs](#) Grafana Products Open Source Learn Downloads Login Contact us

Grafana Cortex Graphite **Loki** Metictank Prometheus Tanka



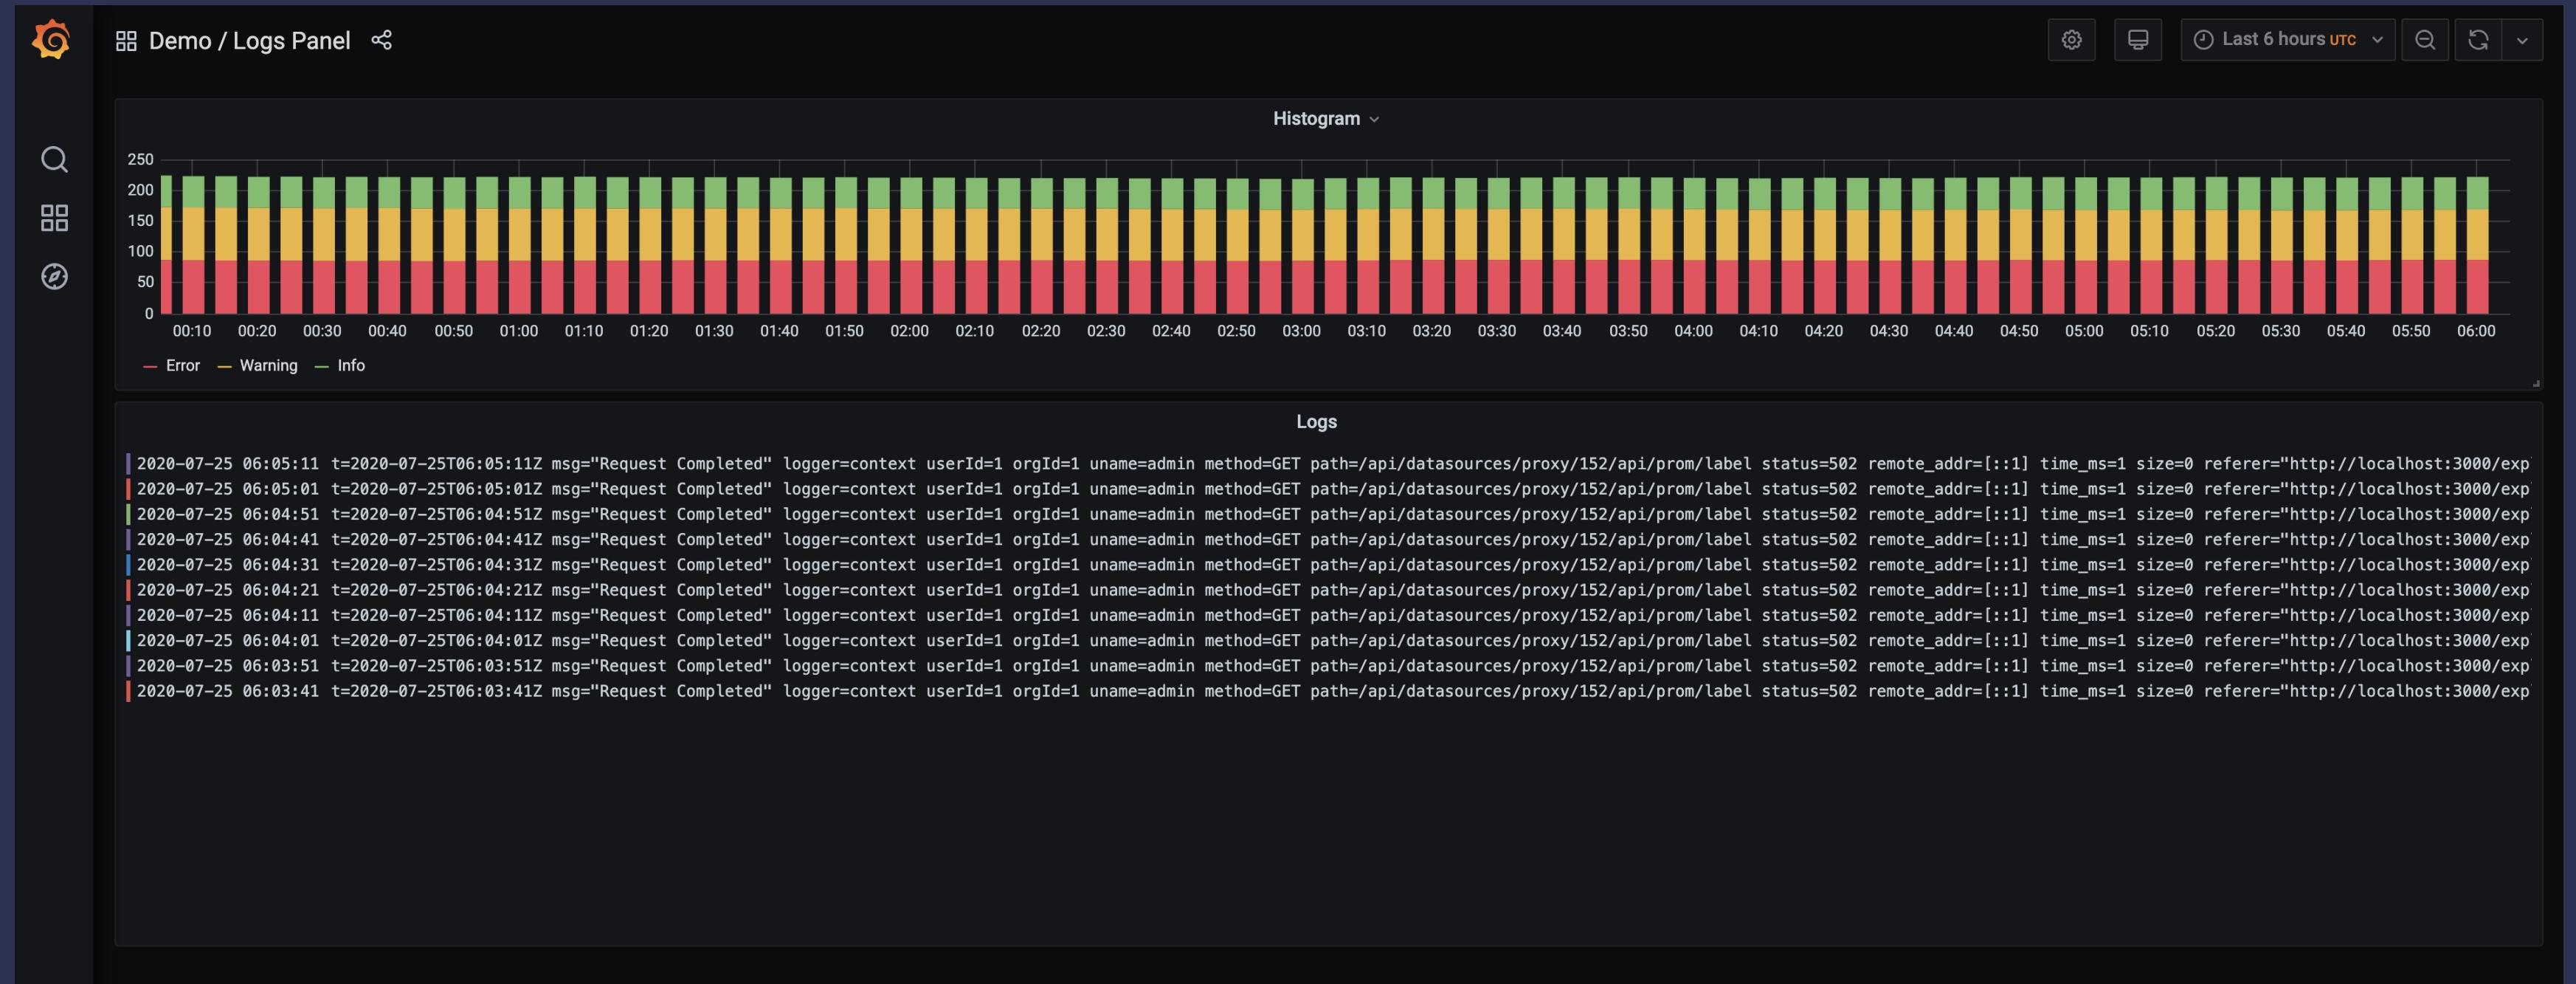
Loki

Loki is a horizontally-scalable, highly-available, multi-tenant log aggregation system inspired by Prometheus. It is designed to be very cost effective and easy to operate. It does not index the contents of the logs, but rather a set of labels for each log stream.

The Loki project was started at Grafana Labs in 2018, and announced at [KubeCon Seattle](#). Loki is released under the Apache 2.0 License.

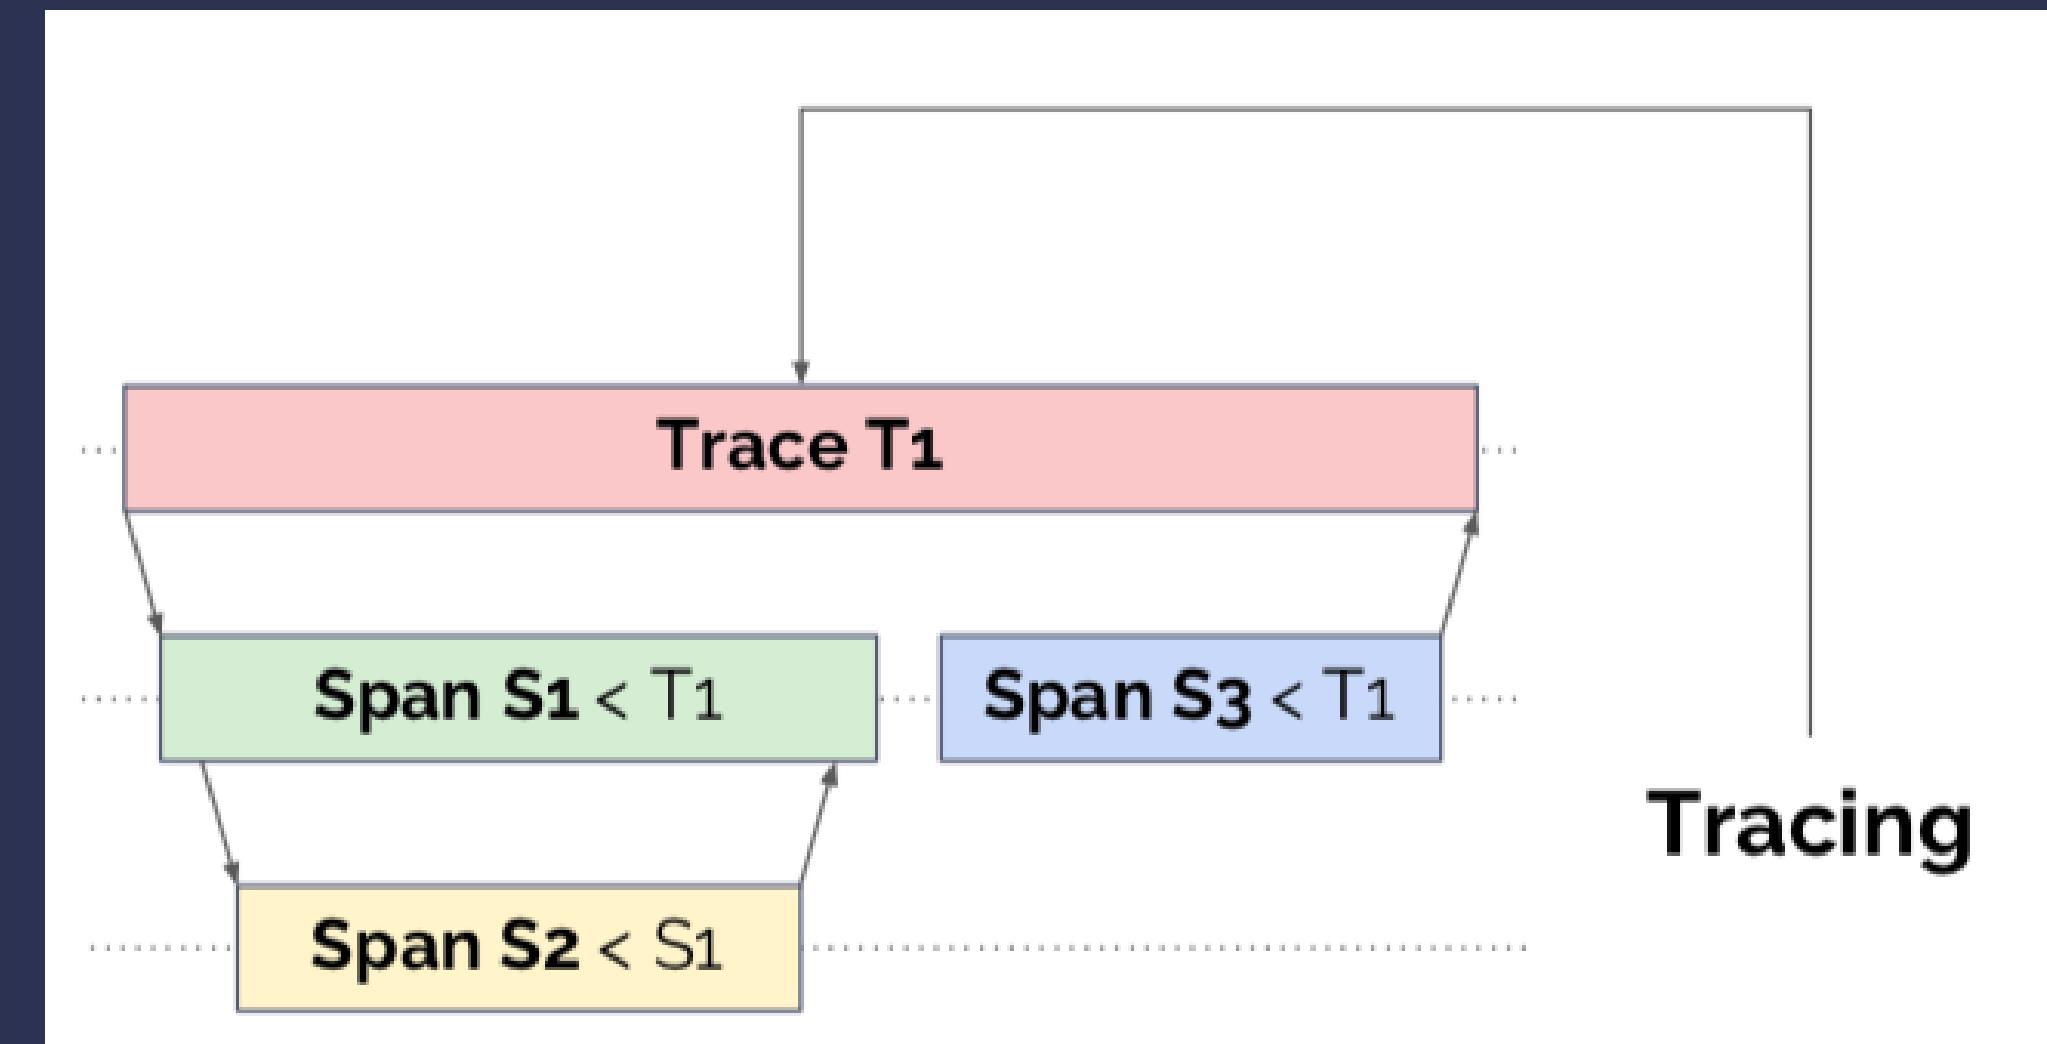
Grafana Labs is proud to lead the development of the Loki project, building first-class support for Loki into Grafana, and ensuring Grafana Labs customers receive Loki support and features they need.

- [Loki Documentation](#)
- [Loki Releases](#)
- [Loki Source](#)



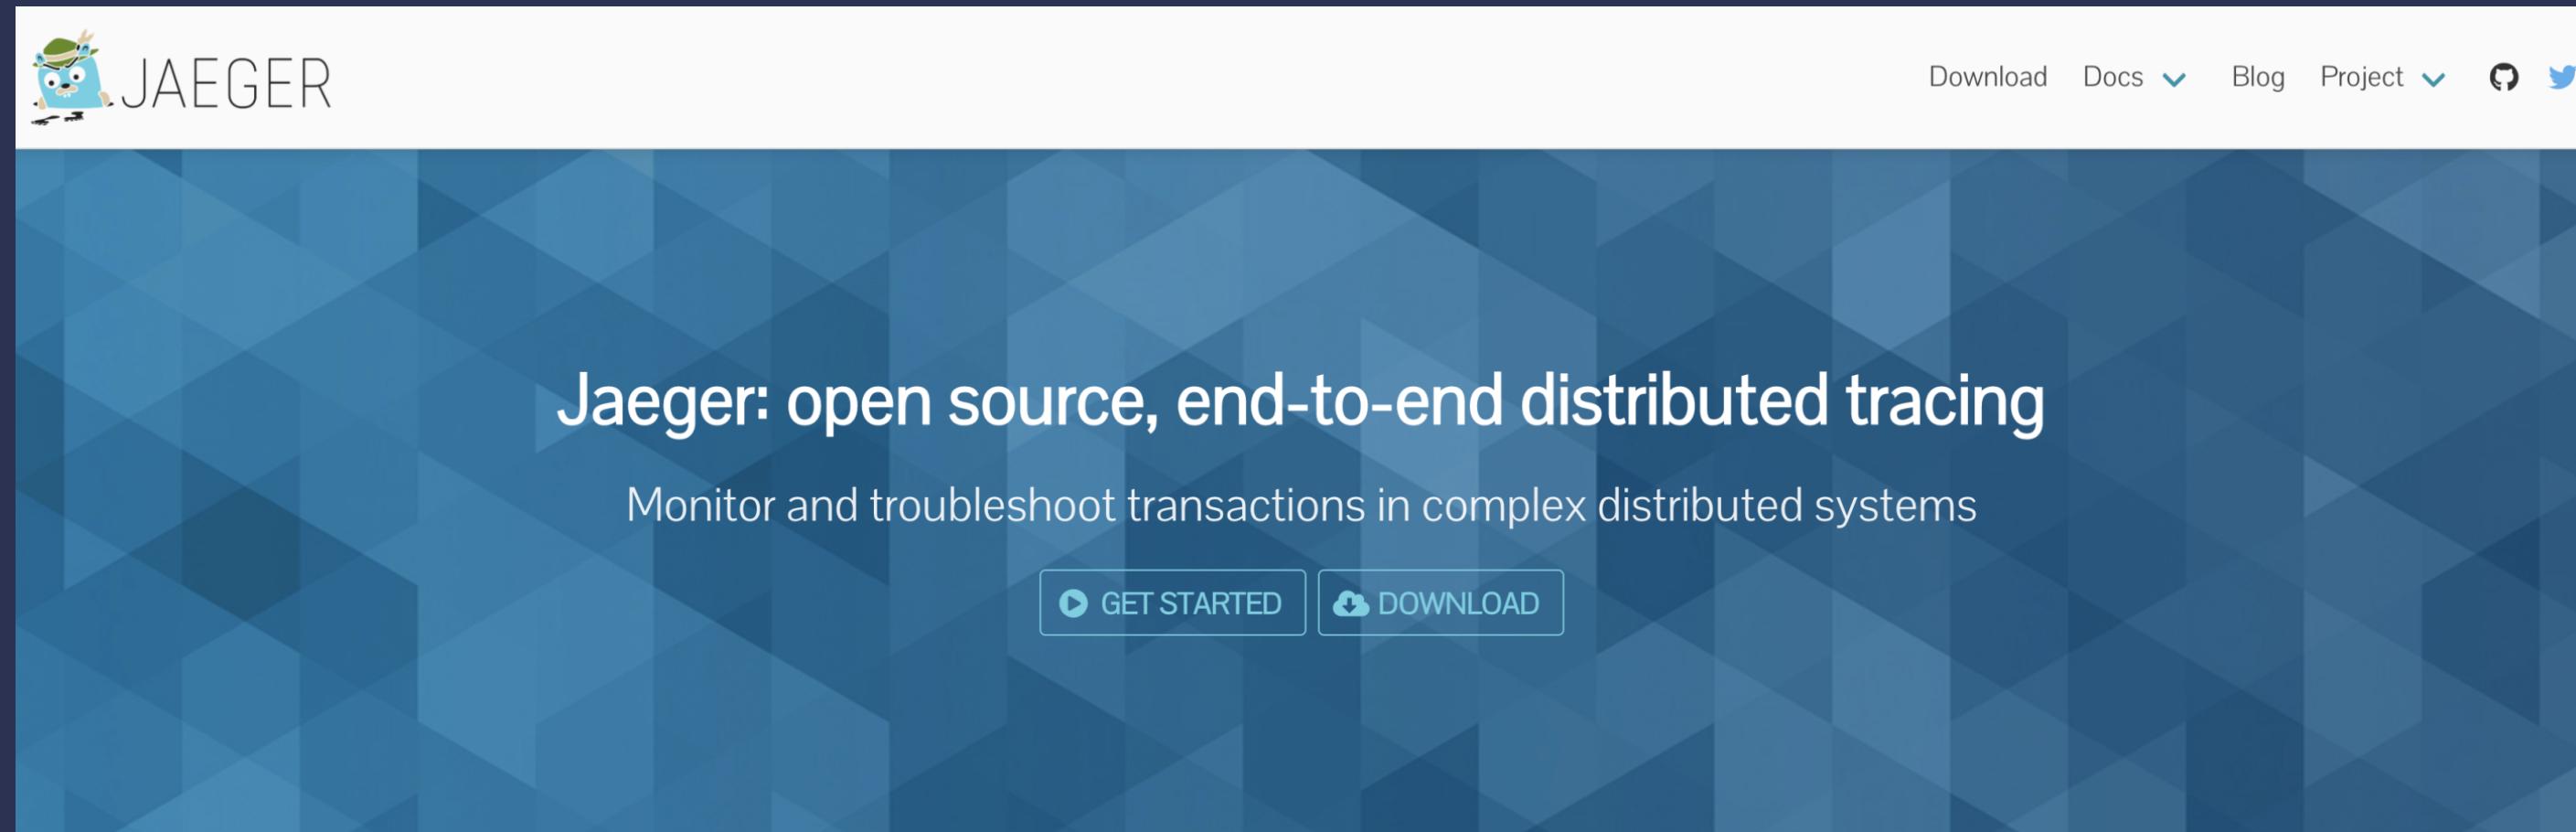


Tracing





jaeger



The screenshot shows the Jaeger homepage. At the top left is the Jaeger logo (an owl icon) and the word "JAEGER". At the top right are links for "Download", "Docs", "Blog", "Project", and social media icons for GitHub and Twitter. The main title "Jaeger: open source, end-to-end distributed tracing" is centered above a subtitle "Monitor and troubleshoot transactions in complex distributed systems". Below the subtitle are two buttons: "GET STARTED" and "DOWNLOAD". The background of the main section is a blue geometric pattern.

Why Jaeger?

As on-the-ground microservice practitioners are quickly realizing, the majority of operational problems that arise when moving to a distributed architecture are ultimately grounded in two areas: **networking** and **observability**. It is simply an orders of magnitude larger problem to network and debug a set of intertwined distributed services versus a single monolithic application.

Problems that Jaeger addresses

				
distributed transaction monitoring	performance and latency optimization	root cause analysis	service dependency analysis	distributed context propagation



```
// Create and install Jaeger export pipeline.  
traceProvider, closer, err := jaeger.NewExportPipeline(  
    jaeger.WithCollectorEndpoint("http://127.0.0.1:14268/api/traces"),  
    jaeger.WithProcess(jaeger.Process{  
        ServiceName: serviceName,  
        Tags: []kv.KeyValue{  
            kv.String("exporter", "jaeger"),  
        },  
    }),  
    jaeger.WithSDK(& sdktrace.Config{DefaultSampler: sdktrace.AlwaysSample  
})  
if err != nil {  
    stdlog.Fatalf("failed to initialize tracer, err: %v", err)  
}  
  
defer closer()
```



```
func Receive(logger log.Logger, tracer trace.Tracer) http.HandlerFunc {
    return func(w http.ResponseWriter, r *http.Request) {
        ctx, span := tracer.Start(r.Context(), "receive")
        defer span.End()

        ...

        var reqBuf []byte

        if err := tracer.WithSpan(ctx, "decode", func(ctx context.Context
            var err error
            reqBuf, err = snappy.Decode(nil, compressed)
            return err
        )); err != nil {
            level.Warn(logger).Log("msg", "snappy decode", "err", err)
            http.Error(w, err.Error(), http.StatusBadRequest)
        }
    }
}
```



```
func Receive(logger log.Logger, tracer trace.Tracer) http.HandlerFunc {
    return func(w http.ResponseWriter, r *http.Request) {
        ctx, span := tracer.Start(r.Context(), "receive")
        defer span.End()

        ...

        var req prompb.WriteRequest

        if err := tracer.WithSpan(ctx, "unmarshal", func(ctx context.Context) {
            return proto.Unmarshal(reqBuf, &req)
        }); err != nil {
            level.Warn(logger).Log("msg", "proto unmarshalling", "err", err)
            http.Error(w, err.Error(), http.StatusBadRequest)

        }
    }
}
```



```
// Tracer returns an HTTP handler that injects the given tracer and start
// If any client span is fetched from the wire, we include that as our pa
func Tracer(logger log.Logger, tracer trace.Tracer, name string) func(next
    operation := fmt.Sprintf("/%s HTTP[server]", name)

    ...
    return http.HandlerFunc(func(w http.ResponseWriter, r *http.Reque
        ctx := r.Context()
        attrs, entries, spanCtx := httptrace.Extract(ctx, r)
        r = r.WithContext(correlation.ContextWithMap(ctx, correlation
            MultiKV: entries,
        })))
        ...
    })
}
}
```



```
// Tracer returns an HTTP handler that injects the given tracer and start
// If any client span is fetched from the wire, we include that as our pa
func Tracer(logger log.Logger, tracer trace.Tracer, name string) func(next
    operation := fmt.Sprintf("/%s HTTP[server]", name)
    ...
    return http.HandlerFunc(func(w http.ResponseWriter, r *http.Reque
        ...
        ctx, span := tracer.Start(
            trace.ContextWithRemoteSpanContext(ctx, spanCtx),
            name,
            trace.WithAttributes(attrs...),
        )
        defer span.End()

        span.AddEvent(ctx, operation)
```



Jaeger UI [Lookup by Trace ID...](#) [Search](#) [Compare](#) [Dependencies](#) [About Jaeger](#)

Search [JSON File](#)

Service (4)

Operation (25)

Tags ?

Lookback

Min Duration

Max Duration

Limit Results

[Find Traces](#)



Logo by Lev Polyakov
www.polyakovproductions.com



Jaeger UI [Lookup by Trace ID...](#) [Search](#) [Compare](#) [Dependencies](#) [About Jaeger](#)

Search [JSON File](#)

Service (4)
thanos-query

Operation (25)
all

Tags (1)
http.status_code=200 error=true

Lookback
Last Hour

Min Duration
e.g. 1.2s, 100ms, 500us

Max Duration
e.g. 1.2s, 100ms, 500us

Limit Results
20

[Find Traces](#)

Duration

Time

08:01:07 am 08:01:08 am 08:01:09 am

20 Traces

Sort: [Most Recent](#) [Deep Dependency Graph](#)

Compare traces by selecting result items

Trace ID	Spans	Timestamp	Duration
48de4eb	40 Spans	Today 8:01:09 am	2.11ms
7ad4f13	46 Spans	Today 8:01:08 am	1.94ms
3d465ed	46 Spans	Today 8:01:08 am	2.06ms

thanوس-query: /query HTTP[server] 48de4eb
40 Spans thanос-query (30) thanос-receive (9) thanос-rule (1)
Today | 8:01:09 am a few seconds ago

thanос-query: /query HTTP[server] 7ad4f13
46 Spans thanос-query (30) thanос-receive (15) thanос-rule (1)
Today | 8:01:08 am a few seconds ago

thanос-query: /query HTTP[server] 3d465ed
46 Spans thanос-query (30) thanос-receive (15) thanос-rule (1)
Today | 8:01:08 am a few seconds ago

52 / 79

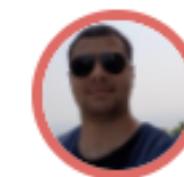




⌚ 14:00 - 15:00

Distributed Tracing on Go using OpenTelemetry

In the microservices era, monitoring applications become harder due to distributed and asynchronous data flow, therefore quickly identifying... [Show More](#)



Fahri Yardımcı



Profiling



```
func main() {
    debug := os.Getenv("DEBUG") != ""
    if debug {
        runtime.SetMutexProfileFraction(cfg.debug.mutexProfileFraction)
        runtime.SetBlockProfileRate(cfg.debug.blockProfileRate)
    }
}
```



```
import "net/http/pprof"

// Internal server to expose introspection APIs.
mux := http.NewServeMux()

// Register pprof endpoints.
mux.HandleFunc("/debug/pprof/", pprof.Index)
mux.HandleFunc("/debug/pprof/cmdline", pprof.Cmdline)
mux.HandleFunc("/debug/pprof/profile", pprof.Profile)
mux.HandleFunc("/debug/pprof/symbol", pprof.Symbol)
mux.HandleFunc("/debug/pprof/trace", pprof.Trace)
```



localhost:8091/debug/pprof/

/debug/pprof/

Types of profiles available:

Count Profile

30	allocs
0	block
0	cmdline
18	goroutine
30	heap
0	mutex
0	profile
18	threadcreate
0	trace

[full goroutine stack dump](#)

Profile Descriptions:

- allocs: A sampling of all past memory allocations
- block: Stack traces that led to blocking on synchronization primitives
- cmdline: The command line invocation of the current program
- goroutine: Stack traces of all current goroutines
- heap: A sampling of memory allocations of live objects. You can specify the gc GET parameter to run GC before taking the heap sample
- mutex: Stack traces of holders of contended mutexes
- profile: CPU profile. You can specify the duration in the seconds GET parameter. After you get the profile file, use the go tool pprof command to view it.
- threadcreate: Stack traces that led to the creation of new OS threads
- trace: A trace of execution of the current program. You can specify the duration in the seconds GET parameter. After you get the trace file, use the go tool pprof command to view it.



Profiling Go with pprof ↗

Julia Evans

ABOUT
TALKS
PROJECTS
TWITTER
GITHUB

FAVORITES ★ ZINES ★ RSS

Profiling Go programs with pprof

Last week me and my cool coworker Josh were debugging some memory problems in a Go program using [pprof](#).

There's a bunch of pprof documentation on the internet but I found a few things confusing so here are some notes so I can find them easily.

First – when working with pprof it's good to be running a recent version of Go! For example Go 1.8 adds [mutex profiles](#) so you can see mutex contention.

in this post I'll

- link to the useful pprof resource I found
- explain what a pprof profile is



Continues Profiling



Continues Profiling Go Programs

Continuous Profiling of Go programs



Jaana Dogan 

Apr 26, 2018 · 4 min read



One of the most interesting parts of Google is our fleet-wide continuous profiling service. We can see who is accountable for CPU and memory usage, we can continuously monitor our production services for contention and blocking profiles, and we can generate analysis and reports and easily can tell what are some highly impactful optimization projects we can work on.

I briefly worked on [Google Cloud Profiler](#), our new product that is filling the cloud-wide profiling gap for Cloud users. Note that you DON'T need to run your code on Google Cloud Platform in order to use it. Actually, I use it



Conprof ↗

Why GitHub? Team Enterprise Explore Marketplace Pricing

Search Sign in Sign up

conprof / conprof

Code Issues 9 Pull requests 3 Actions Security Insights

master 2 branches 0 tags Go to file Code

brancz Merge pull request #66 from yeya24/add-profile-download-api ... 3d56f13 11 days ago 125 commits

File	Description	Time Ago
api	Add Profile Type	2 months ago
config	Add Profile Type	2 months ago
deployments	fixes after review	last month
examples	examples: Configure default config to scrape conprof itself	2 years ago
internal/trace	Add make sync	2 months ago
pkg/runutil	*: Fix versioning	12 months ago
pprofui	support downloading profiles	11 days ago
scrape	Add Profile Type	2 months ago
scripts	*: Fix versioning	12 months ago
vendor	*: Use go 1.13	9 months ago
web	Add Profile Type	2 months ago
.dockerignore	Add container make target	14 months ago
.drone.yml	Add arm build to latest	3 months ago
.gitattributes	Add generated git attribute	2 months ago
.gitignore	deployments: Rework to be more idiomatic jsonnet structure	9 months ago
.golangci.yml	*: Fix versioning	12 months ago
.promu.yml	*: Use go 1.13	9 months ago
Dockerfile	Add deployments/ dir containing kubernetes manifests	14 months ago
LICENSE	Initial commit	2 years ago

About

Continuous profiling in for pprof compatible profiles.

pprof profiling continuous-profiling
profiles prometheus

Readme Apache-2.0 License

Releases No releases published

Contributors 11

Languages

Go 69.6% TypeScript 15.6%
Jsonnet 6.4% JavaScript 4.0%
Makefile 2.2% HTML 1.8%
Other 0.4%



Conprof

From {job="conprof", profile_path="/debug/pprof/heap"} 04/23/2019, 07:06 PM To 04/23/2019, 08:26 PM Now >

{__address__="localhost:8080",__scheme__="http",instance="localhost:8080",job="conprof",profile_path="/debug/pprof/heap"}



2019/4/23 20:33

2019/4/23 20:40

2019/4/23 20:46

2019/4/23 20:53

2019/4/23 21:00



Conprof

{container="thanos-receive",profile_path="/debug/pprof/heap"}

From

14.07.2020, 10:43

To

25.07.2020, 05:09



Now



{__address__="10.128.12.130:10902",__scheme__="http",container="thanos-receive",instance="10.128.12.130:10902",job="thanos",namespace="telemeter-production",pod="observatorium-thanos-receive-default-5",profile_path="/debug/pprof/heap"}

2020/7/23 15:13

2020/7/24 00:56

2020/7/24 10:40

2020/7/24 17:15

2020/7/25 06:00

{__address__="10.128.2.173:10902",__scheme__="http",container="thanos-receive",instance="10.128.2.173:10902",job="thanos",namespace="telemeter-production",pod="observatorium-thanos-receive-default-0",profile_path="/debug/pprof/heap"}

2020/7/23 15:13

2020/7/24 00:56

2020/7/24 10:40

2020/7/24 20:23

2020/7/25 06:00

{__address__="10.128.4.81:10902",__scheme__="http",container="thanos-receive",instance="10.128.4.81:10902",job="thanos",namespace="telemeter-production",pod="observatorium-thanos-receive-default-4",profile_path="/debug/pprof/heap"}

2020/7/23 15:13

2020/7/24 00:56

2020/7/24 10:40

2020/7/24 20:23

2020/7/25 06:00

{__address__="10.128.9.236:10902",__scheme__="http",container="thanos-receive",instance="10.128.9.236:10902",job="thanos",namespace="telemeter-production",pod="observatorium-thanos-receive-default-2",profile_path="/debug/pprof/heap"}

2020/7/23 15:13

2020/7/24 00:56

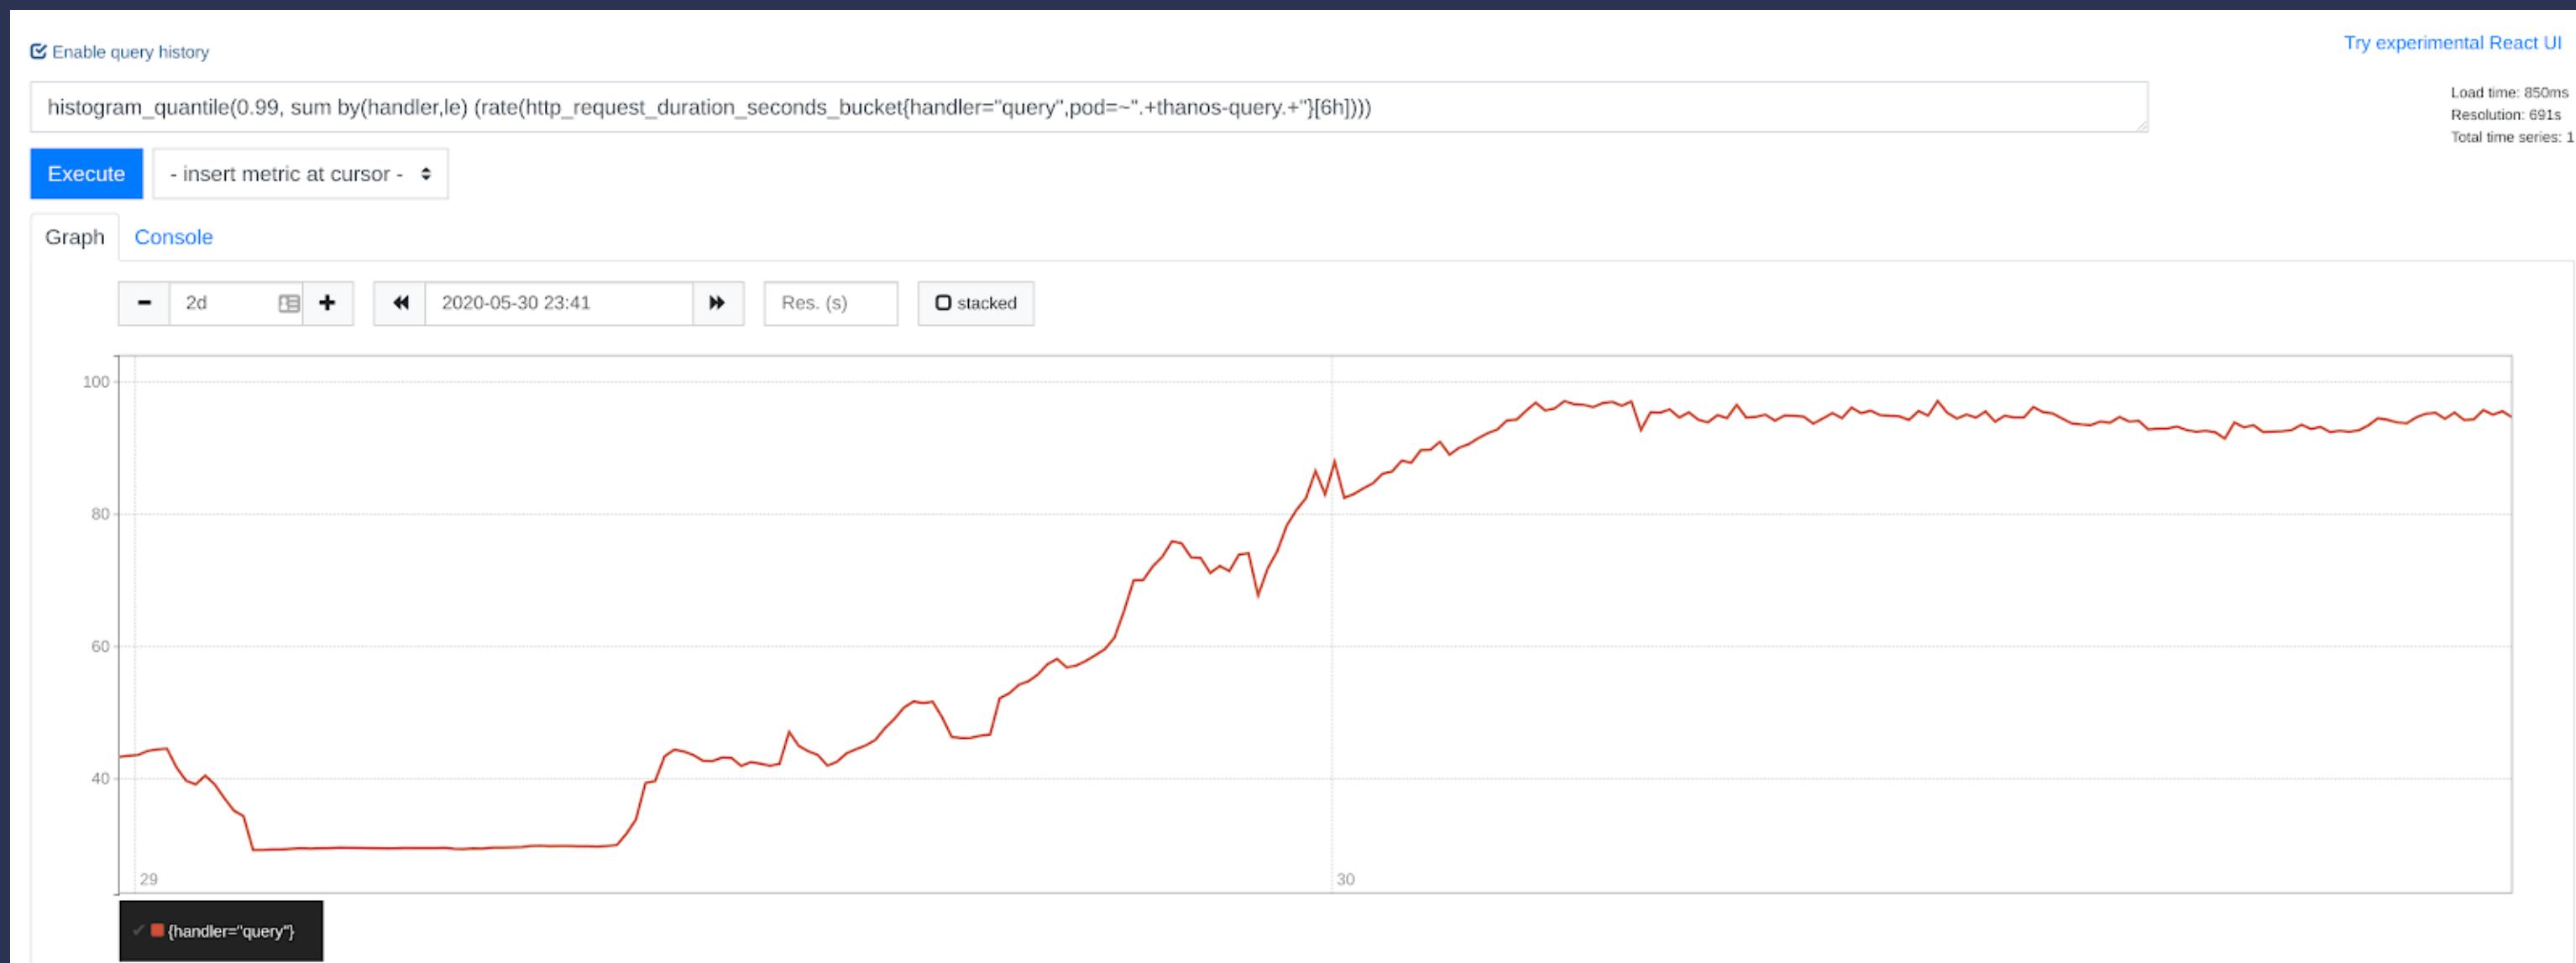
2020/7/24 10:40

2020/7/24 20:23

2020/7/25 06:00

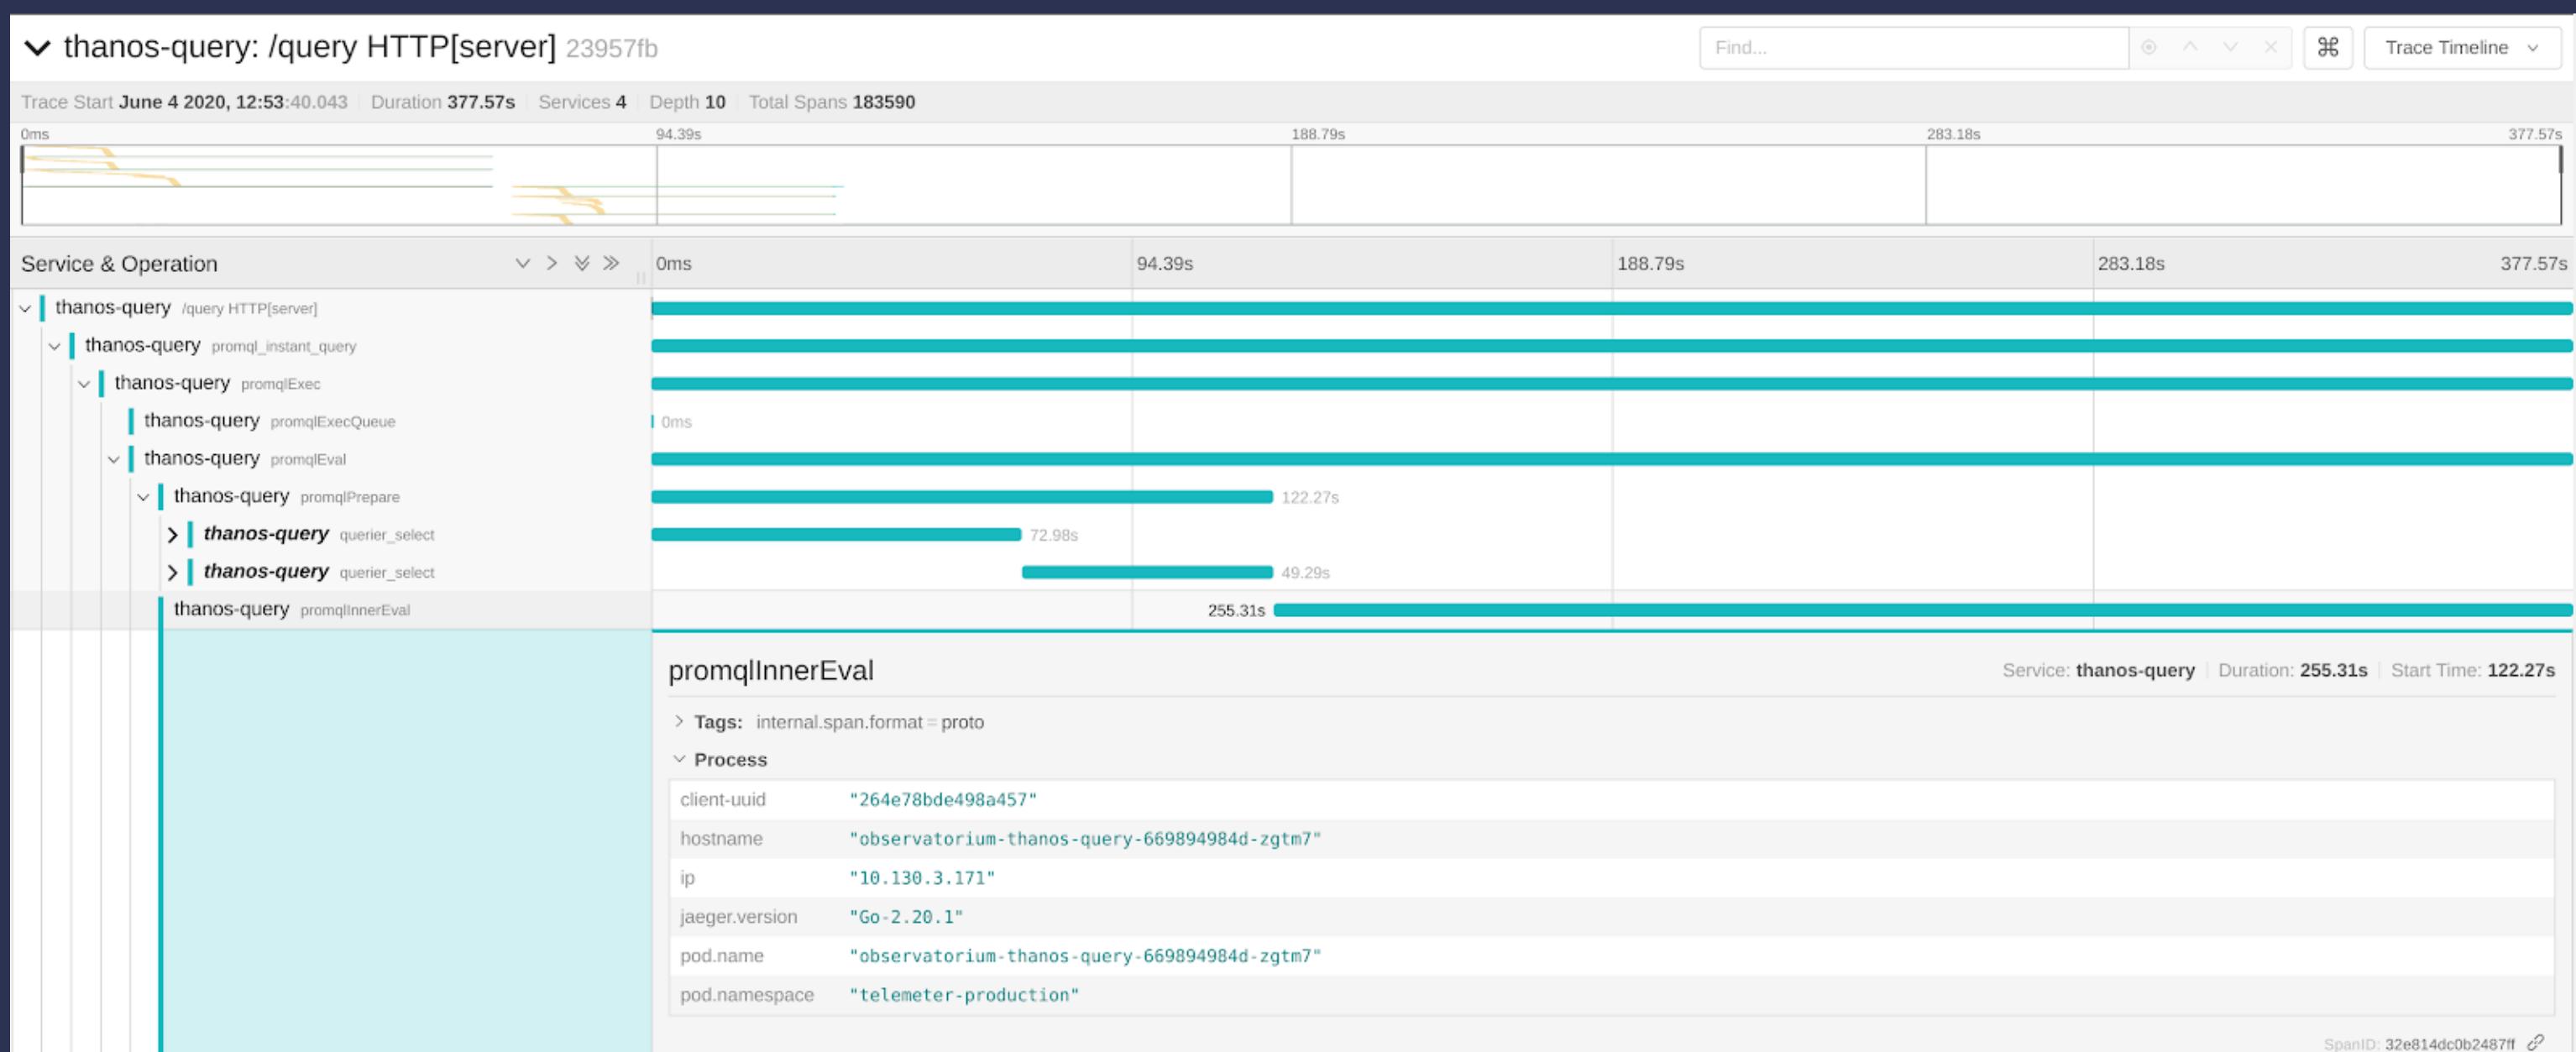


Correlation: Observability Superpower





```
ts=2020-06-04T10:59:57.6170538Z component=query-reader msg="request finis  
Log line from a slow query log, showing a trace ID from a slow query.
```





Conprof

From 06/04/2020, 11:04 AM To 06/04/2020, 01:05 PM Now >

```
{__address__="10.130.3.171:9090",__scheme__="http",container="thanos-query",instance="10.130.3.171:9090",job="thanos",namespace="telemeter-production",pod="observatorium-thanos-query-669894984d-zgtm7",profile_path="/debug/pprof/profile"}
```

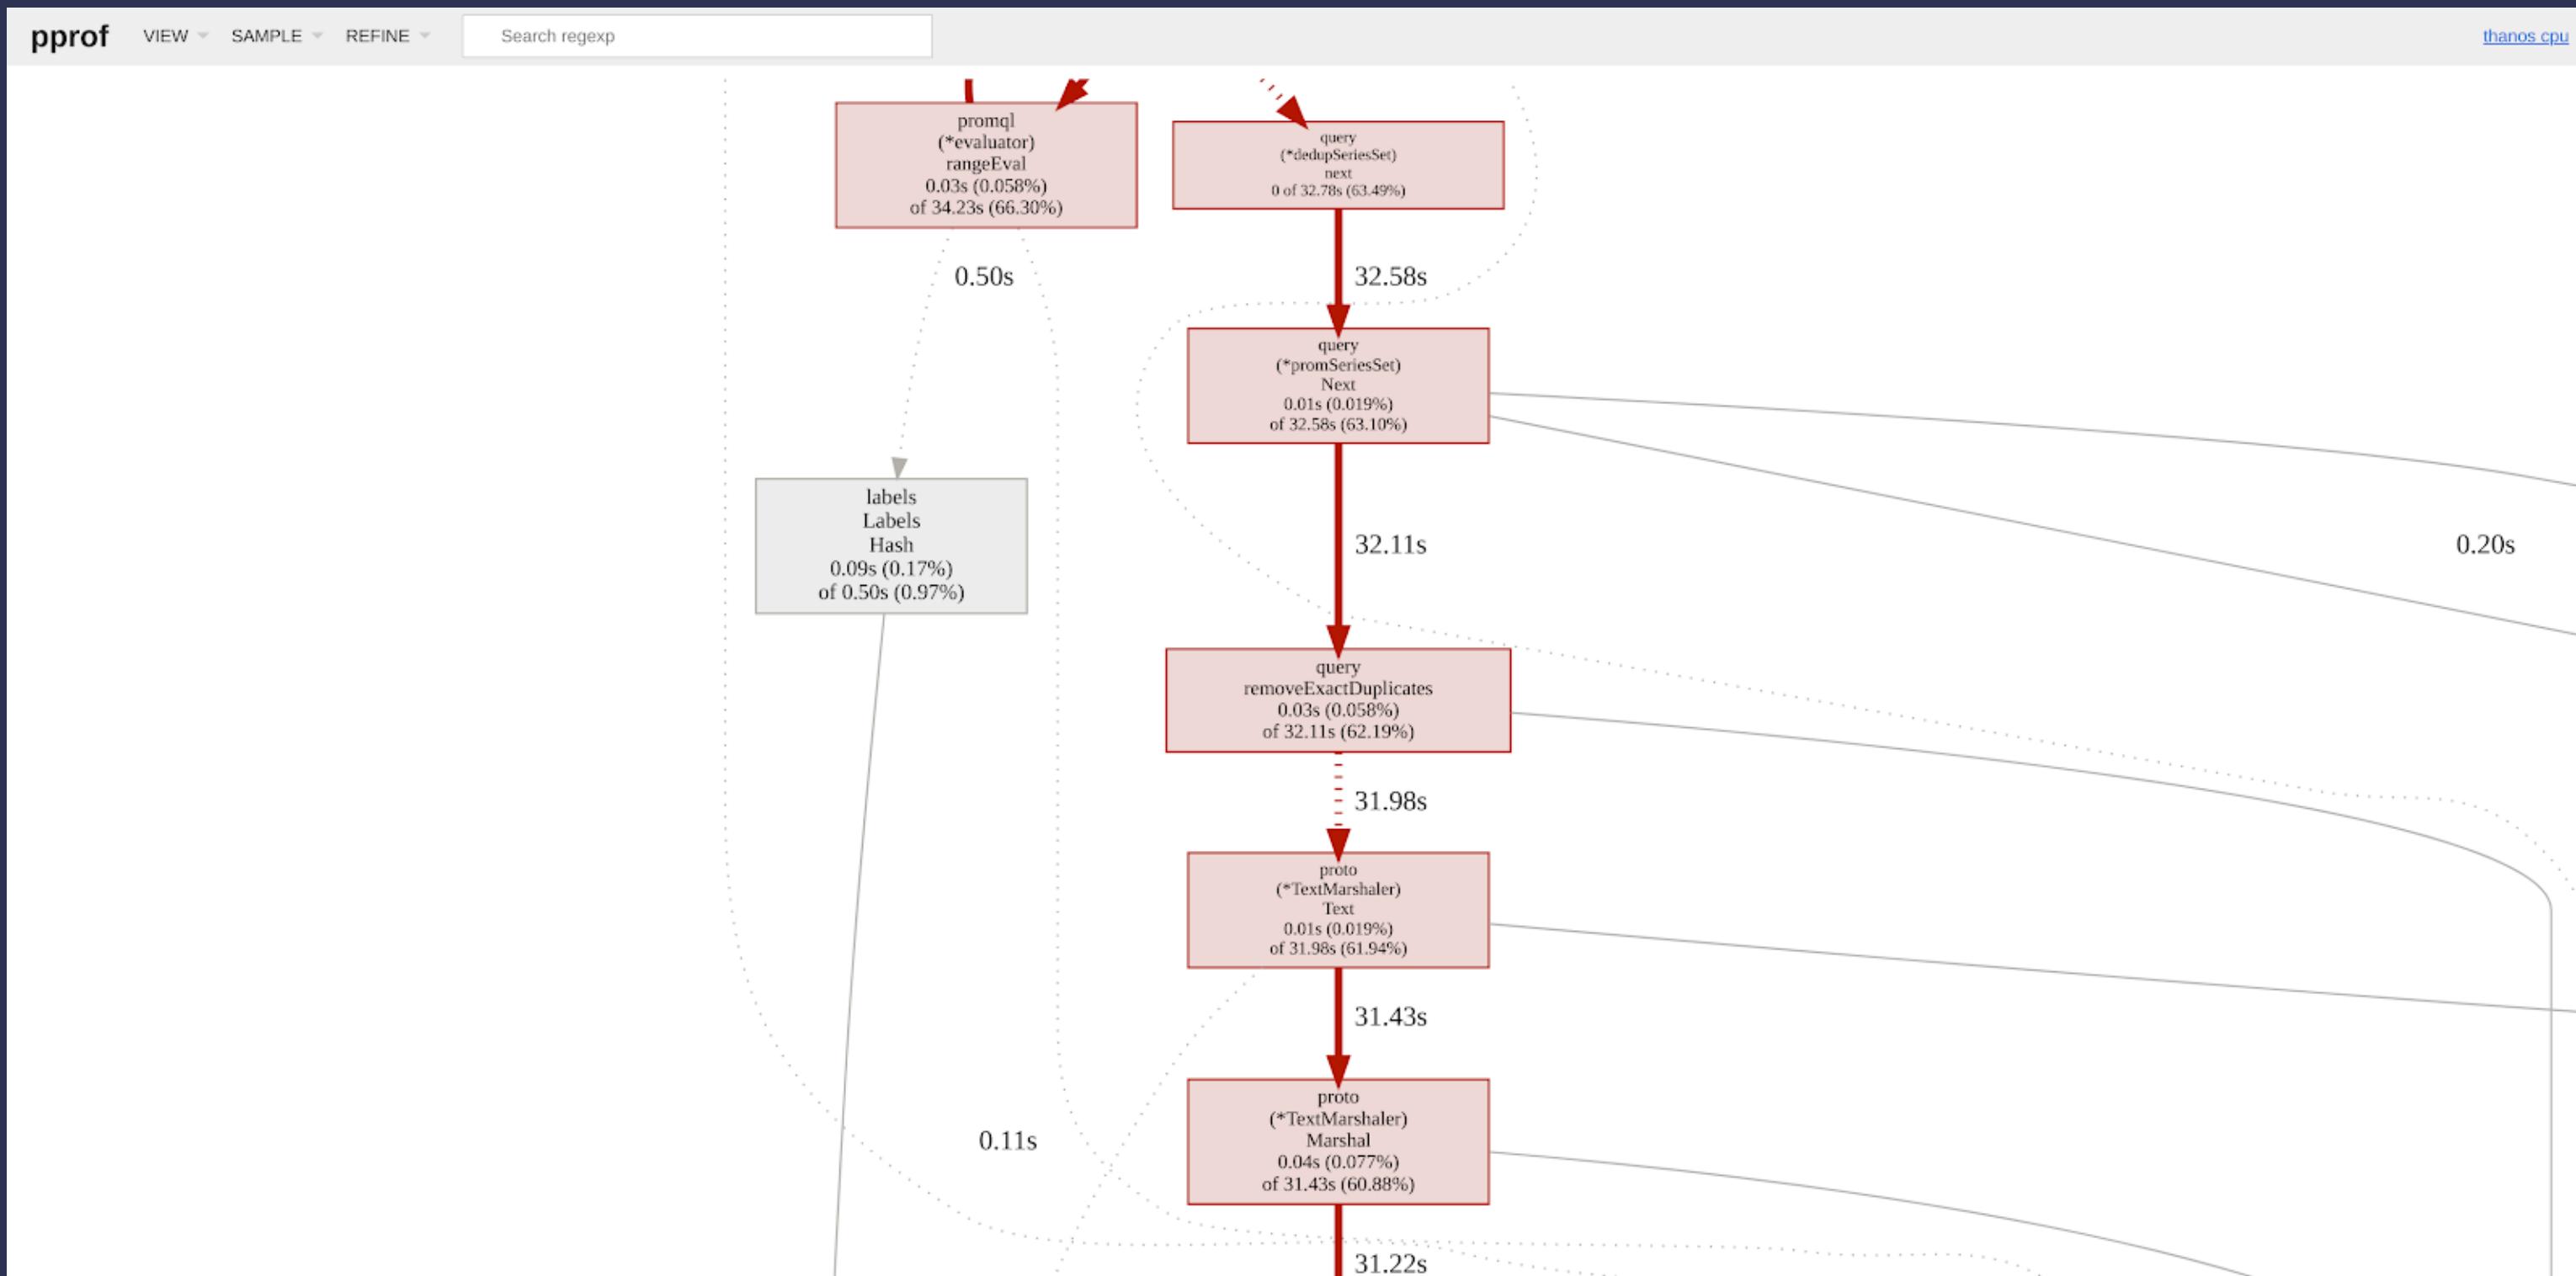
2020/6/4 10:40

2020/6/4 11:30

2020/6/4 12:20

2020/6/4 13:10

2020/6/4 14:01





Grafana



Grafana



<https://play.grafana.org/d/000000100/mixed-datasources>



Prometheus data source

 [Data sources](#) > Prometheus

Grafana includes built-in support for Prometheus. This topic explains options, variables, querying, and other options specific to the Prometheus data source. Refer to [Add a data source](#) for instructions on how to add a data source to Grafana.

Prometheus settings

To access Prometheus settings, click the **Configuration** (gear) icon, then click **Data Sources**, and then click **Prometheus**.

Name	Description
Name	The data source name. This is how you refer to the data source in panels and queries.
Default	Default data source means that it will be pre-selected for new panels.
Url	The URL of your Prometheus server, e.g. <code>http://prometheus.example.org:9090</code> .
Access	Server (default) = URL needs to be accessible from the Grafana backend/server, Browser = URL needs to be accessible from the browser.
Basic Auth	Enable basic authentication to the Prometheus data source.
User	User name for basic authentication.
Password	Password for basic authentication.
Scrape interval	Set this to the typical scrape and evaluation interval configured in Prometheus. Defaults to 15s.



Using Loki in Grafana

› Data sources › Loki

BETA: Querying Loki data requires Grafana's Explore section. Grafana v6.x comes with Explore enabled by default. In Grafana v5.3.x and v5.4.x. you need to enable Explore manually. Viewing Loki data in dashboard panels is supported in Grafana v6.4+.

Grafana ships with built-in support for Loki, Grafana's log aggregation system. Just add it as a data source and you are ready to query your log data in [Explore](#).

Adding the data source

1. Open Grafana and make sure you are logged in.
2. In the side menu under the `Configuration` link you should find a link named `Data Sources` .
3. Click the `Add data source` button at the top.
4. Select `Loki` from the list of data sources.

Note: If you're not seeing the `Data Sources` link in your side menu it means that your current user does not have the `Admin` role for the current organization.

Name	Description
Name	The data source name. This is how you refer to the data source in panels, queries, and Explore.
Default	Default data source means that it will be pre-selected for new panels.



jaeger-datasource ✅



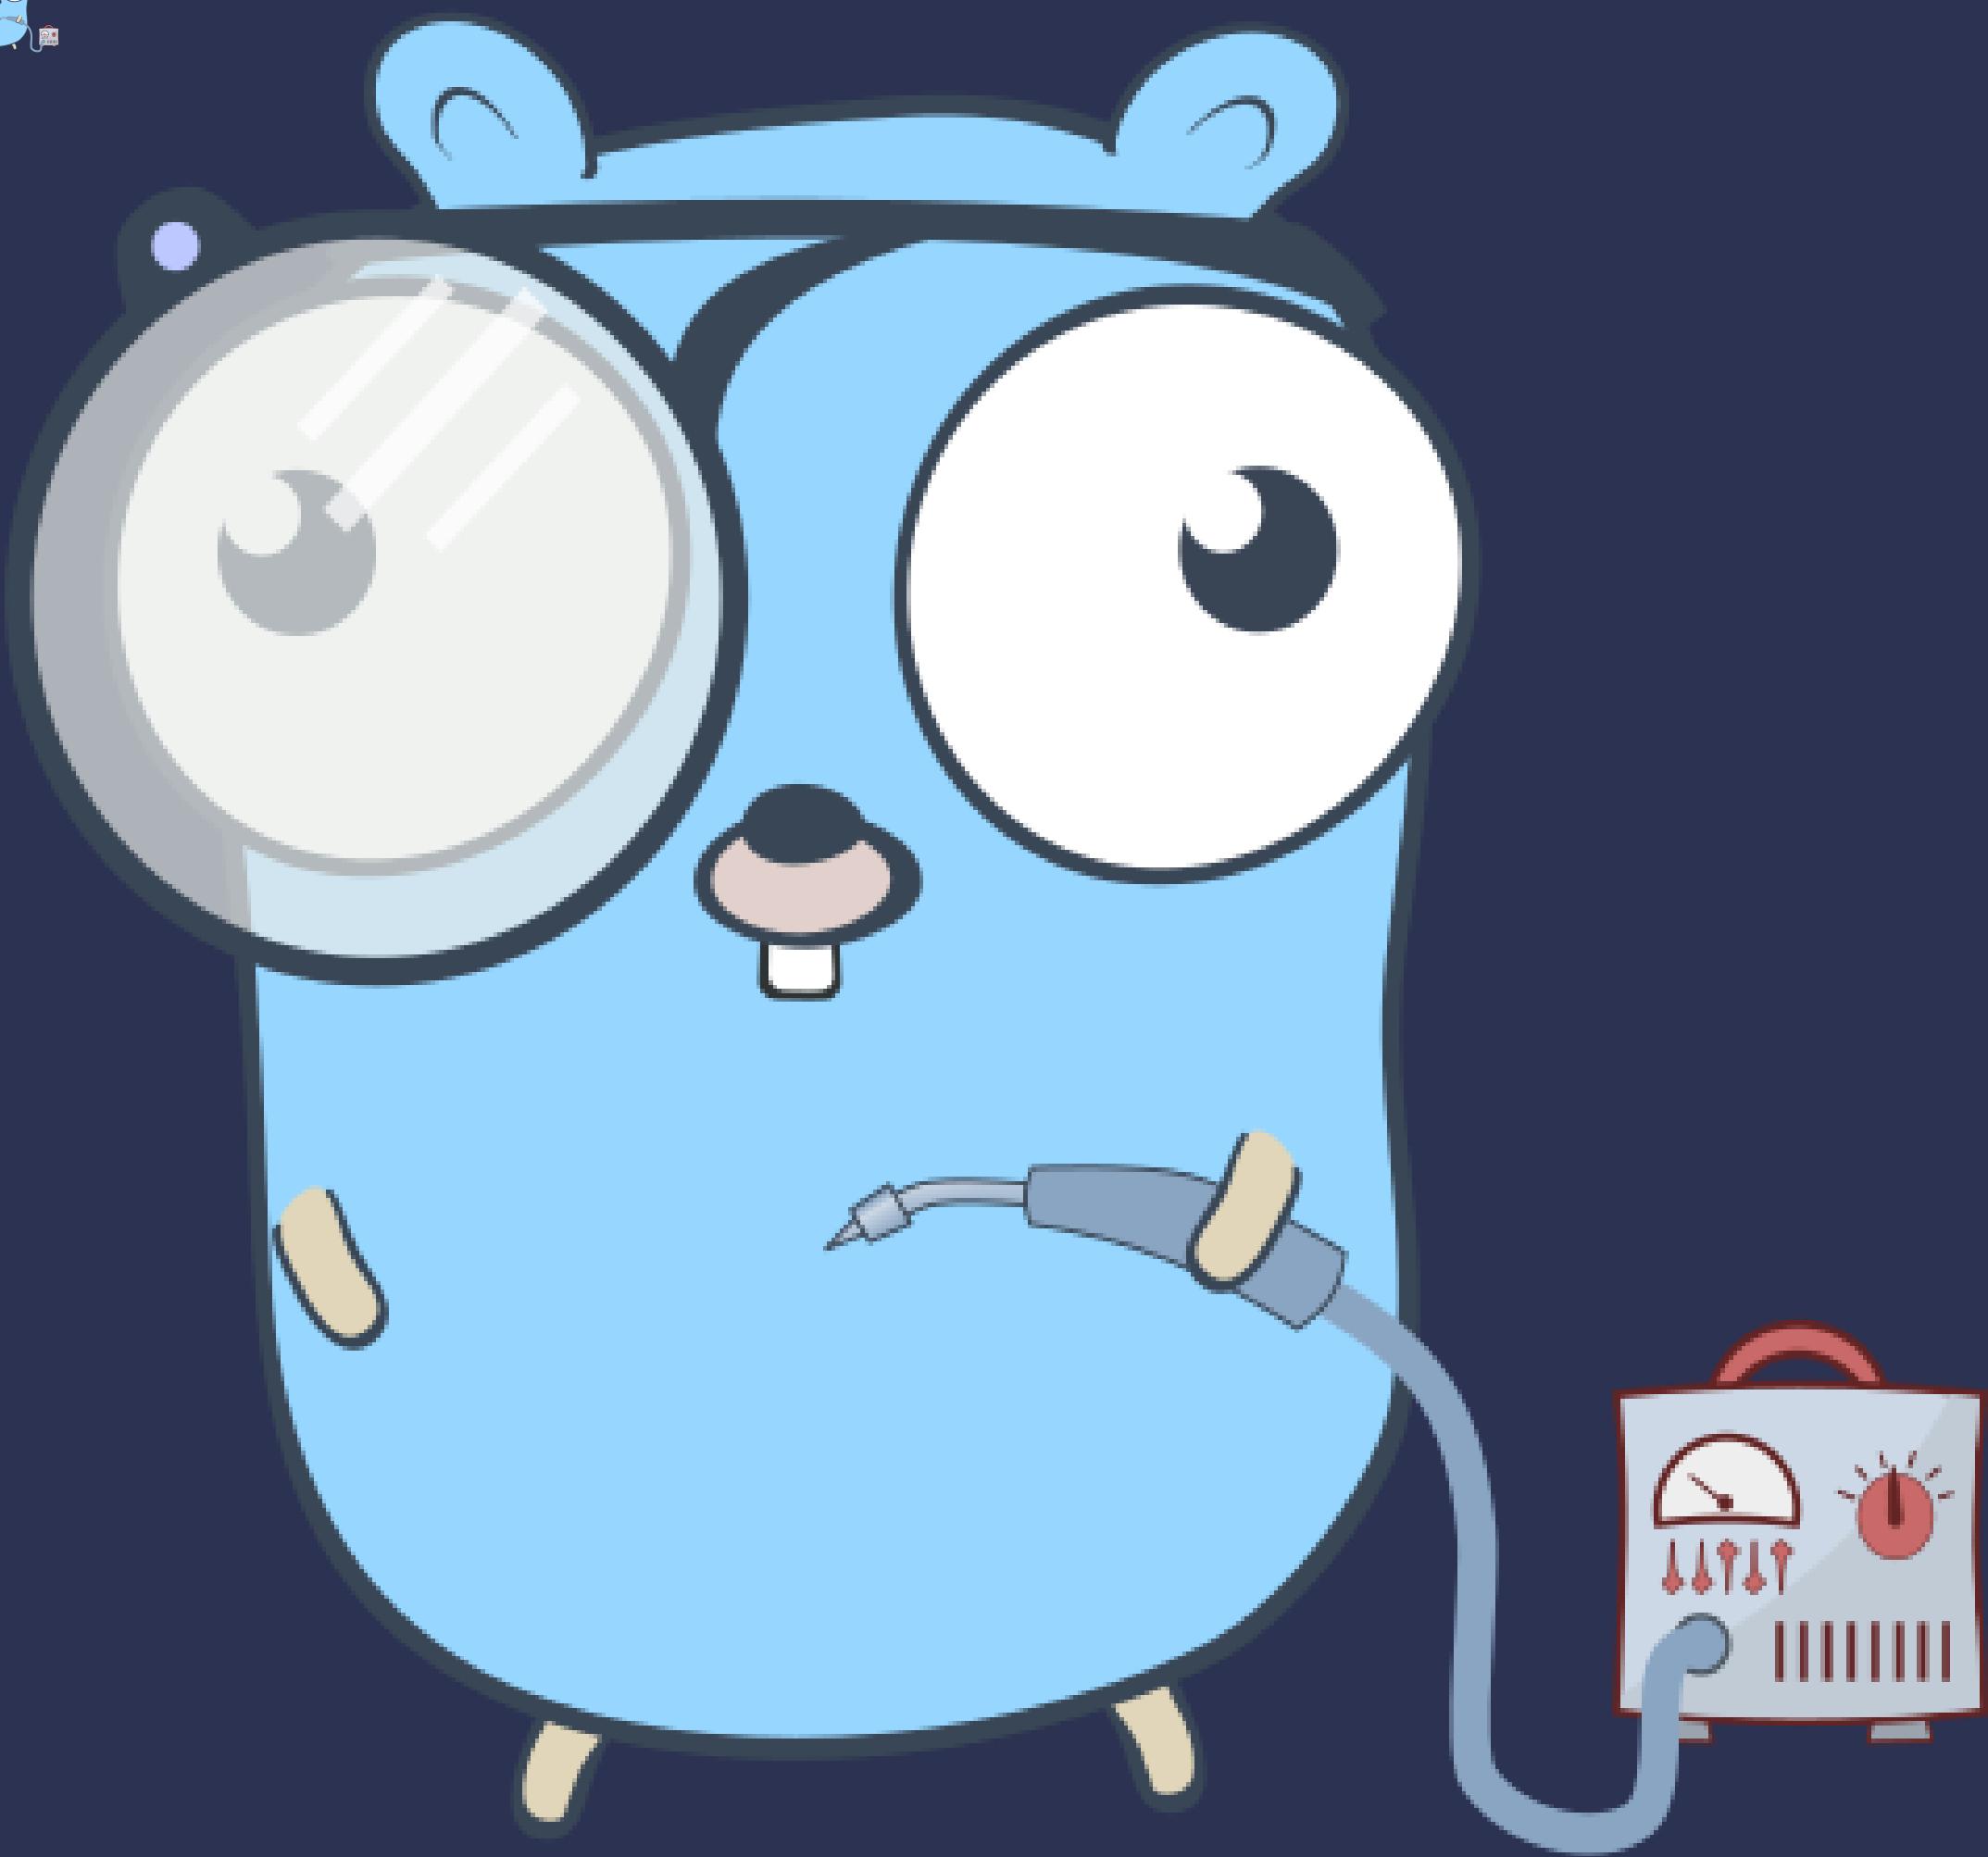
Soon Conprof datasource

conprof-datasource 



Check it yourself

observable-remote-write ✎



Thanks a lot





Red Hat